Accuracy Is Not the Objective: A Cross-Disciplinary Review and Global Study of Surrogate Pipelines for Expensive Engineering Analyses

Working draft — review part

J. Poznanski

October 31, 2026

Abstract

Draft. Surrogate models stand in for expensive simulations inside design loops. Engineering practice is dominated by Gaussian-process (kriging) surrogates, while tabular machine learning and several simulation-heavy sciences are dominated by tree ensembles; a thesis finite-element study found gradient-boosted trees out-predicting every Gaussian-process variant on real data while losing on smooth synthetic surfaces. This paper reviews surrogate, uncertainty, sampling, multi-fidelity, change-of-basis, upscaling and surrogate-based-optimization methods across disciplines, from computation theory to engineering practice, and then (in the study part, in preparation) scores complete pipelines by the true-function regret they deliver per unit of simulation cost.

1 Introduction

Design and analysis loops—optimization, uncertainty propagation, calibration, reliability—call a simulation many thousands of times. When one call costs minutes to hours, as for the coupled finite-element analyses behind the present work, the loop is run on a *surrogate*: a cheap model fitted to a limited set of simulation results [1–3]. The idea has been reinvented under many names in many fields—response surfaces in chemical engineering [4], kriging in mining geostatistics [5–7], design and analysis of computer experiments in statistics [8, 9], proxy models in petroleum engineering [10], emulators in climate science, cosmology and epidemiology [11–13], metamodels in simulation and operations research [14], surrogate benchmarks in automated machine learning [15], and fitness approximation in evolutionary computation [16, 17].

Each community has converged on its own default. In engineering design the Gaussian process (GP) is the dominant surrogate [2, 18], and multi-fidelity GP formulations are the natural extension when cheaper analyses are available [19–21]. In tabular machine learning, tree ensembles remain state of the art at small-to-medium sample sizes [22–24], and tree-based surrogates are preferred for algorithm-configuration benchmarks [15, 25] and agent-based model calibration [26]. These defaults are rarely confronted with each other on the same problems, with the same budgets, under the same noise.

The confrontation is not academic. The surrogate study of Poznański [27, Sec. 14.3 and App. G], which sized the deployable duct wall of a tube-launched ramjet from coupled fluid–structure finite-element analyses, reported a GP with the lowest error on six synthetic surfaces, yet gradient-boosted trees with the lowest cross-validated error on the real finite-element data; the trees were set aside because the downstream interior-point optimizer requires smooth derivatives. That study selected its

surrogate by approximation accuracy, implicitly assuming that the most accurate surrogate yields the best design—an assumption contradicted by several benchmarks in which the approximation winner and the optimization winner differ [15, 28, 29]. The present paper revisits and extends that study; where its findings are corrected here, the corrections are stated explicitly.

This paper has two parts. The **review** (Sections 2–9) surveys, breadth first and across disciplines, the methods that make up a surrogate pipeline: the learner (Section 2), its uncertainty estimate and the sampling it drives (Section 3), multi-fidelity fusion (Section 4), change-of-basis and latent-space representations borrowed from data-driven reduced-order modelling (Section 5), the conversion of a learner into an optimizer-ready representation (Section 6), and surrogate-based optimization itself (Section 8), followed by the theory that predicts when each should win (Section 9). We deliberately include old methods that were not pursued and methods that are standard outside aerospace but rare within it. The **study** (in preparation) varies surface class, dimension, budget, noise level and type, fidelity structure and cost, scout–learner pairing, upscaling and optimizer type in a factorial design, and scores pipelines by true-function regret per unit simulation cost. Section 11 states the gap the study fills.

2 The landscape of surrogate models

2.1 Lineages

Surrogate families descend from five traditions, and each inherited its home field's assumptions.

Geostatistics → kriging → Gaussian processes. Krige's regression corrections for gold-grade estimation on the Witwatersrand [5] were formalised by Matheron as optimal linear prediction of a random function [6]; Cressie traces the distinction between the two [7]. Sacks, Welch, Mitchell and Wynn transplanted the method to deterministic computer codes, treating the code output as a realisation of a stochastic process [8], which founded the design and analysis of computer experiments [9, 14, 18]. The same predictor is the GP regression of machine learning [30]. Extensions address trends (universal and blind kriging [31]), gradients [32], high dimension through partial least squares [33], input-dependent and stochastic noise [34, 35], scale [36, 37], and non-stationarity by warping, partitioning or depth [38–40]. The treed GP was itself motivated by a rocket-booster CFD experiment [39].

Approximation theory → splines, radial basis functions, sparse grids, rational functions. Hardy's multiquadric for topographic surfaces [41] and Duchon's thin-plate splines [42] began scattered-data approximation, now a mature theory with native-space error bounds [43–45]; compactly supported kernels make it sparse [46]. Tensor-product B-splines [47] give smooth, cheap, differentiable representations on structured grids, at a cost exponential in dimension; Smolyak's construction [48, 49] breaks that cost for functions of mixed smoothness. Chebyshev technology [50] and rational approximation, revived by the AAA algorithm [51], offer spectral accuracy for analytic functions and natural representation of pole-like features such as the $1/R$ collapse knee of our application.

Statistics → regression, splines and additive models. Response-surface methodology [4] fits low-order polynomials; sparse regression [52] selects terms from large dictionaries, as in ALAMO [53] and, for dynamics, SINDy [54]. Smoothing splines are the GP posterior mean for a particular prior [55, 56]; penalised B-splines (P-splines) fit splines directly to scattered, noisy data [57];

additive models and thin-plate regression splines scale them up [58, 59]; shape constraints encode monotonicity [60, 61]. Friedman's MARS [62] builds adaptive hinge bases that represent kinks exactly, and projection pursuit regression [63] sums ridge functions of learned projections—an early learned change of basis (Section 5).

Machine learning → trees, ensembles, neural networks, foundation models. Regression trees [64] partition the input space into boxes; bagging them gives random forests [65, 66], boosting them gives gradient-boosting machines [67] and their engineered descendants [68–70]; Bayesian sums of trees give BART [71], which has been used with sequential design for computer experiments [72]. Linear-leaf model trees [73] and Mondrian forests [74] are less common variants. Neural surrogates range from shallow perceptrons to deep ensembles [75], Bayesian networks [76–78], derivative-aware Sobolev training [79], Kolmogorov–Arnold networks with spline edges [80] and, most recently, tabular foundation models that perform Bayesian inference in context [81]. Physics-informed networks and neural operators [82–84] target field outputs and are outside our scalar-response scope. The polynomial neural networks of Ivakhnenko's group method of data handling [85]—self-organising, validation-selected polynomial networks from the late 1960s—are a largely abandoned branch of the same tree.

Physics and numerical analysis → spectral expansions and mapped models. Polynomial chaos [86, 87] expands the response in orthogonal polynomials of the inputs; sparse variants recover it from few runs [88]. Space mapping [89, 90] aligns a cheap physics model with an expensive one rather than fitting data from scratch (Section 4).

Simple local methods—Shepard's inverse-distance weighting [91], Sibson's natural-neighbour interpolation [92], moving least squares [93], LOESS [94] and Lipschitz interpolation [95]—are mostly absent from engineering comparisons but are the right baselines: as the climate-emulation literature shows, simple methods can match or beat large models when noise is significant [96].

2.2 What each field optimises for

The disagreement between fields is partly explained by what they benchmark on. Engineering comparisons [97–103] typically use smooth analytic functions, small samples, and deterministic outputs; GPs and RBFs win there, as they did in the synthetic-surface study of [27, App. G]. Williams and Cremaschi trained 16 632 models of eight families on 99 functions of the Surjanovic–Bingham library across five shape classes and found the winner depended on shape: GP most often on valley, bowl and miscellaneous shapes, MARS on many-local-minima functions; the best single predictor of accuracy was a spacing statistic of the design, not a landscape feature [28, 104]. In tabular machine learning, Grinsztajn et al. isolated three mechanisms behind the trees' advantage: neural networks are biased toward overly smooth functions (smoothing the training target hurt trees and barely affected networks), are harmed more by uninformative inputs, and are rotation-invariant where the data are not [22, 105, 106]. A stationary-kernel GP shares the first bias and, with an isotropic kernel, the third. In agent-based modelling, ragged simulator responses led Lamperti et al. to boosted trees precisely because kriging's smoothness assumptions "collapse interesting patterns" [26]. Process systems engineering, water resources, building energy, reliability, reservoir engineering, finance, turbomachinery and additive manufacturing each report their own mix [10, 107–116].

The common lesson of every large benchmark is that there is no universal winner but there are predictable ones [29, 88, 97, 104, 112, 117]. Two further lessons shape our design. Within-family configuration can matter as much as the choice of family—solver and sampler choices changed sparse-PCE errors by orders of magnitude [88]—so each family must be given its best configuration

under an equal tuning budget. And ensembles of surrogates weighted by cross-validation error [118–120] hedge against picking the wrong family when the predictors of the right one are unavailable.

3 Uncertainty estimates and the sampling they drive

3.1 Uncertainty estimators by family

Adaptive sampling needs a surrogate that knows where it is wrong. GPs provide a posterior variance in closed form [30]; note, however, that for fixed hyperparameters this variance depends only on the input locations, not on the observed outputs, so a greedy maximum-variance rule is a space-filling design in disguise until hyperparameters are re-estimated. Other families need constructed estimates: bootstrap committees and query by committee [121]; the jackknife and infinitesimal jackknife for random forests [122]; quantile regression forests [123]; distributional boosting [124] and virtual ensembles for boosting [125]; Bayesian sums of trees [71]; deep ensembles [75], Monte Carlo dropout [126], Laplace approximations [78] and evidential heads [127, 128] for networks [129]; and model-agnostic conformal intervals [130]. Tabular foundation models return predictive distributions directly [81]. These estimators differ in calibration and, importantly for sampling, in *shape*: a tree committee's spread is piecewise constant, so its maximiser is degenerate over whole cells.

3.2 Initial designs

Space-filling designs—Latin hypercubes [131], maximin designs [132], low-discrepancy sequences [133, 134] and centroidal Voronoi tessellations [135]—bound integration error through discrepancy and approximation error through fill distance [43]. Several large studies found the choice among them second-order: Williams and Cremaschi saw no significant difference between Halton, Sobol' and LHS [104], and in high dimension no advanced scheme beat LHS for sparse PCE [88]. The more consequential comparison is space-filling versus *independent random* sampling, which is the control most adaptive-sampling studies report, including [27, App. G].

3.3 Adaptive sampling for global accuracy

Reviews of adaptive sampling for global metamodelling [136–138] and a replicable comparison for kriging [117] catalogue the main criteria: maximum variance and active learning [139]; integrated variance reduction [140]; mutual information [141]; expected improvement for global fit [142]; expected prediction error [143]; local-nonlinearity-weighted Voronoi exploration [144]; Taylor-expansion error indicators [145]; and cross-validation-based criteria [146]. Fuhg et al. conclude that success depends on the problem and the goal of the analysis [117]. Moustapha et al. make the same point for reliability, where 39 combinations of surrogate, reliability estimator, learning function and stopping rule were run on 20 problems: no strategy dominated, but feature-dependent recommendations emerged [112]. Their modular decomposition is the template for our stage-swap ablation.

3.4 Sampling for a goal

When the goal is an optimum rather than a global fit, acquisition functions trade exploration for exploitation: probability of improvement [147], expected improvement and EGO [148], confidence bounds with regret guarantees [149], knowledge gradient [150], entropy search variants [151, 152],

constrained variants [153], trust-region schemes for high dimension [154], random-forest and density-ratio surrogates [155, 156], and stochastic RBF methods [157, 158]. When the goal is a boundary, as for reliability or for the feasible/ infeasible and floor-bound regions of the duct-sizing problem of [27, Sec. 14.3], contour-targeting criteria apply [159, 160]. A pipeline that samples for global accuracy and then optimizes—the pipeline of [27, Sec. 14.3]— therefore spends part of its budget where the optimizer will never go; whether that matters is an empirical question we test.

4 Multi-fidelity modelling

Most engineering analyses come with cheaper approximations: coarser meshes, lower-order physics, closed-form sizing rules. Peherstorfer, Willcox and Gunzburger classify methods that exploit them as *adaptation*, *fusion* and *filtering*, and stress keeping the high-fidelity model in the loop for guarantees [161, 162]; other reviews organise the surrogate side [163, 164].

Correction (bridge) functions. The oldest approach scales or shifts the low-fidelity model: $f_H \approx \beta(x) f_L + \delta(x)$, with β from sensitivities [165], kriging [166] or trust-region model management that enforces first-order consistency [167, 168]. The regression target of [27, Sec. 14.3, App. G]—the ratio of finite-element to closed-form thickness—is exactly a multiplicative bridge, an unnamed multi-fidelity model.

Autoregressive co-kriging. Kennedy and O'Hagan model each level as a scaled previous level plus an independent GP discrepancy [19]; Forrester, Sóbester and Keane brought it to optimization [20]; Le Gratiet and Garnier showed a recursive formulation gives the same predictor with independent, cheaper fits [146]; Han and Görtz's hierarchical kriging uses the low-fidelity predictor as the trend [169].

Nonlinear fusion. When fidelities are related nonlinearly, NARGP feeds the low-fidelity prediction into the high-fidelity kernel [170], deep GPs generalise it [171], and composite neural networks learn linear and nonlinear correlations in parallel [172]. Space mapping aligns inputs instead of outputs [89, 90]. For tree ensembles and generic learners, the simplest fusion—adding the low-fidelity prediction as an input feature—is natural but, to our knowledge, rarely benchmarked against co-kriging.

When it helps. Toal's guidance for multi-fidelity kriging is that fidelities must be strongly correlated (squared correlation around 0.9) and that roughly 10–80% of the budget should go to low fidelity [173]. Van Rijn and Schmitt show that correlation alone is a poor guide and propose "error grids" over the numbers of high- and low-fidelity samples, estimated by subsampling a single design [174], using the benchmark collection of [175]. The NATO AVT-354 group proposes analytical benchmarks spanning localised, multimodal, discontinuous and noisy behaviour together with global-accuracy and goal-oriented metrics [176, 177]; cost-aware multi-fidelity acquisition extends Bayesian optimization [178]. In hypersonic vehicle design, multi-fidelity GP surrogates with trajectory-informed and adjoint-guided sampling make coupled aero-thermal-trajectory optimization tractable [21, 179].

What is missing is a comparison of multi-fidelity methods *across learner families*, under matched cost, as a function of fidelity correlation, relationship type, cost ratio and budget split—and on a real problem where the low-fidelity model is available everywhere, as the closed-form sizing rule is in our finite-element data.

5 Changing the domain: latent, sparse-basis and low-rank surrogates

Data-driven reduced-order modelling succeeds by finding a basis in which dynamics are simple: proper orthogonal decomposition and dynamic mode decomposition find low-rank linear structure, SINDy finds a sparse dictionary representation [54], and the SINDy autoencoder learns the coordinates and the sparse model jointly [180]. Compressed and randomised DMD do not remove any stage; they compute the same decomposition from a sketch of the data, with cost set by the intrinsic rank rather than the ambient size [181–183]. The same ideas apply to static response surfaces, where they are scattered across communities that rarely cite each other.

Linear input reduction. Projection pursuit regression sums smooth functions of learned projections [63]; active subspaces take the dominant eigenvectors of the expected gradient outer product [184]; polynomial ridge approximation fits the subspace and the polynomial jointly without gradients [185]; GPs with built-in dimension reduction learn the projection inside the kernel [186]; random embeddings exploit low intrinsic dimension for optimization [187].

Co-learned nonlinear maps. Deep kernel learning places a GP on the output of a neural feature map and trains both by marginal likelihood [188]—the most direct form of "learning the domain and the fit together". Its weakness is decisive at engineering sample sizes: with many hyperparameters the marginal likelihood prefers to correlate all points, and the method can over-fit worse than a plain network unless treated fully Bayesianly [189]. Kernel flows learn kernels by a cross-validation-like criterion instead [190]; input warping restricts the map to per-axis monotone transforms [38]; Fourier features and spectral-mixture kernels learn spectral content [191, 192]. Latent-space optimization with variational autoencoders works for structured inputs [193], but an autoencoder trained only on uniformly sampled inputs learns the sampling density, not the response; for static surrogates the map must be supervised.

Sparse bases and compressed sensing. If the response is sparse in a known basis, ℓ_1 recovery needs samples in proportion to the sparsity rather than the basis size [194, 195]; for orthogonal polynomial bases this yields compressive-sampling polynomial chaos [196–198], and wavelets do the same for localised features [199]. The guarantees need random samples from the right measure, which couples the basis to the sampling design.

Low rank: the singular value decomposition of a surface. A bivariate function admits a Schmidt (functional SVD) expansion, the basis of two-dimensional Chebyshev technology [200]; in higher dimension the tensor-train format stores a function with cost linear in dimension and learns it by cross approximation from adaptively chosen samples [201, 202], with continuous analogues for surrogates [203]. Randomised Tucker and tensor-train decompositions compress structured grids [204]; Nyström and random-feature approximations compress kernel matrices [205, 206]; sketching compresses least squares [207]. These are the static counterparts of compressed DMD: they accelerate an intermediate stage—for us, densification of the intermediate model onto a grid, and repeated refits of the scout—without changing what is computed.

Automatic co-search. Choosing a transform, a basis and a learner together is an instance of the combined algorithm selection and hyperparameter optimization problem of automated machine learning [208, 209]; its selection bias must be controlled by nested validation.

6 From learner to optimizer-ready surrogate

A gradient-based optimizer such as an interior-point method [210] needs a surrogate that is cheap to evaluate and has well-behaved first and second derivatives. Many accurate learners fail one requirement: a GP's prediction costs $O(n)$ per point and its derivatives must be assembled; trees and forests are piecewise constant with zero gradient almost everywhere. A common engineering remedy, and the production pipeline of [27, Sec. 14.3], is a two-stage construction: fit an intermediate model to scattered, adaptively placed samples, densify it onto a structured grid, and interpolate the grid with tensor-product B-splines [47]. Poznański [27, App. G] found the grid–spline step preserved GP accuracy and cut evaluation cost 19–32-fold, but tested a single grid resolution and interpolant.

The literature offers several alternatives with different error and cost profiles. Smoothing splines and P-splines can be fitted directly to scattered noisy data, removing the intermediate model [56, 57, 59]; multilevel B-spline approximation builds a C^2 surface from scattered data by coarse-to-fine lattice refinement [211]; hierarchical and truncated hierarchical B-splines refine locally without a full tensor grid [212, 213]; Smolyak sparse grids reduce the grid size for mixed-smooth functions [48, 49]; Chebyshev interpolation is spectrally accurate for analytic intermediates [50]; low-rank compression of the densified grid (Section 5) reduces storage and evaluation cost from exponential to linear in dimension [201, 204]. The error of any two-stage construction decomposes into the intermediate model's error plus the upscaling error; the second vanishes at the spline's approximation order only when the intermediate model is smooth. For a piecewise-constant intermediate such as a tree ensemble, refining the grid reproduces the staircase more faithfully—which may help accuracy and hurt a gradient-based optimizer at the same time, whereas a *smoothing* upscaler acts as a regulariser. Whether upscaling can turn the most accurate but non-differentiable learner into the best optimizer-facing surrogate is, to our knowledge, untested.

7 Ensemble and stochastic intermediate models

Section 6 treats the intermediate model as one deterministic function to be smoothed. A second line of work keeps *many* models, or a random one, and lets the optimizer see the spread between them. Three uses can be separated: averaging to reduce error, randomness to drive exploration, and stochastic gradients that let first-order methods leave poor basins. They are rarely studied together, and almost never as a choice of optimizer-facing surrogate.

Averaging ensembles. In engineering design the ensemble of surrogates is a weighted average of different model classes, with weights from cross-validation error: the PRESS-weighted average surrogate [118], weights optimised to minimise the ensemble's cross-validation error [214], and the finding that most of the benefit comes from generating many diverse surrogates and selecting by PRESS, with weighting gains that shrink in high dimension [119]. A recent review [215] reports ensembles as the usual remedy for not knowing which surrogate is best, which is the selection problem of this paper seen from the averaging side. In machine learning the ensemble is homogeneous and randomised: bagging fits one class to bootstrap samples [216], and subsampling without replacement (subagging) has the same variance-reduction mechanism with a cleaner analysis [217]; random forests add feature subsampling [65]; deep ensembles train networks from different initialisations [75]; Monte-Carlo dropout samples sub-networks [126]; and randomised prior functions add a fixed random network to each member so that uncertainty exists away from the data, which plain bootstrapping lacks [218]. All of these give a mean predictor and a spread; the spread is what acquisition rules (Section 3) consume.

Randomness for exploration. The member-to-member variation can replace a posterior. Thompson sampling optimises one random draw of the model per step; with ensembles the draw is one member, and the same idea underlies trust-region Thompson sampling in TuRBO [154]. Multiple-surrogate EGO adds one point per surrogate per cycle and imports uncertainty estimates from one surrogate to another, so that models without native uncertainty can be used [219]; bootstrapped kriging replaces the plug-in kriging variance in expected improvement [220]; ensembles of Gaussian processes with different kernels are combined in a mixture posterior for Thompson sampling [221]; and ensemble Bayesian optimization partitions the input space with random Mondrian trees and fits a GP per cell [222]. Neural ensembles are a common alternative to GPs in Bayesian optimization [223], but a systematic comparison found deep ensembles explore poorly and level off below other surrogates on standard benchmarks [224]. Tree ensembles are the surrogate of SMAC [155]. In evolutionary computation, surrogate-assisted memetic algorithms use ensembles and smoothed surrogates inside local search [17, 225], and the stochastic RBF method chooses each new point from random candidates scored by RBF prediction and distance, with probabilistic convergence to the global minimum [157]. For noisy simulators, stochastic kriging separates simulation noise from response-surface uncertainty [35], which is the regime where an averaged intermediate model is most defensible.

Stochastic gradients and smoothing. A deterministic gradient method stops in the first basin. Two mechanisms change that without a new surrogate. First, noise in the gradient: Langevin dynamics adds Gaussian noise to stochastic gradient steps and has non-asymptotic guarantees of reaching approximate global minimisers of non-convex objectives [226]; using one random ensemble member's gradient per step is a model-uncertainty version of the same noise. Second, smoothing: convolving a non-smooth function with a Gaussian or uniform kernel gives a smooth surrogate whose gradient can be estimated from random perturbations [227, 228]. For a piecewise-constant tree ensemble this is the only way to obtain useful gradient information, and a smoothing upscaler (Section 6) is its deterministic counterpart. Particle methods take the ensemble one step further and make the optimizer itself an ensemble: consensus-based optimization moves particles towards a weighted consensus point with an exploration noise [229] and converges globally under mild conditions [230], and ensemble Kalman inversion is a derivative-free ensemble optimizer from inverse problems [231].

Sums of local models. A surrogate that is a sum of simple local pieces is both smooth and naturally stochastic. The partition-of-unity method blends local approximations with weights that sum to one [232]; SNOBFIT fits local quadratic models around the current best points and branches the domain to keep a global search [233]. Drawing one local piece by its weight gives a stochastic gradient whose expectation is the weighted local gradient. We found no study that uses such a decomposition to run stochastic first-order methods on a fitted surrogate.

Optimising tree ensembles directly. Instead of smoothing a tree ensemble, it can be optimised exactly: a tree ensemble is piecewise constant on a finite partition, so its global minimum over a box is a mixed-integer linear program [234]. ENTMOOT solves this with a distance-based uncertainty term for Bayesian optimization over gradient-boosted trees [235], and tree-ensemble kernels give trees a GP-like posterior whose acquisition is a mixed-integer second-order cone program [236]. This is the strongest alternative to upscaling for the thesis' question of whether an accurate tree model can be used for design optimization.

State of the art and gaps. Ensembles are well established for accuracy (averaging) and for exploration (Thompson sampling, multiple-surrogate EGO, SMAC), and tree ensembles can be optimised exactly. What is missing is a matched comparison of these routes from the optimizer's side: (i) averaging versus sampling members versus smoothing versus exact mixed-integer optimization for the same fitted ensemble, scored by regret on the true function rather than by accuracy; (ii) whether the gradient noise from sampling members helps stochastic first-order methods find better optima than the deterministic mean gradient, and at what cost in evaluations; (iii) whether the answer depends on the landscape type, as the accuracy results of this paper suggest it will.

This study covers part of that space. Averaging ensembles (PRESS-weighted average, deep ensembles, Monte-Carlo dropout, random and quantile forests, GBT committees), surrogate-managed loops with randomised acquisition (TuRBO, DYCORS), and two stochastic upscalers with a stochastic-gradient arm (a partition of unity of local quadratics, and a subsample ensemble of smooth members) are included. Methods the literature suggests adding, in order of how directly they answer (i)–(iii): exact mixed-integer optimization of tree ensembles [234, 235]; Thompson sampling over ensemble members as a fit-once optimizer (optimise one random member per restart); Gaussian-smoothing gradient estimates on any surrogate [227]; Langevin-type noise on the mean gradient as the control for member sampling [226]; consensus-based optimization as an ensemble optimizer [229]; optimised-weight ensembles [214] and randomised-prior members [218] as intermediate models; and bootstrapped-kriging expected improvement in the managed loops [220].

8 Surrogate-based optimization and the accuracy fallacy

Optimizers. Once a surrogate exists, any optimizer can be run on it: gradient-based interior-point and SQP methods [210, 237, 238]; derivative-free local methods [239–241]; deterministic global search [242]; and stochastic global methods [243–246]. Surrogate-managed methods instead return to the true function: EGO [148], trust-region model management with provable convergence [167], SNOBFIT [233], stochastic-RBF methods [157, 158], SMAC and TPE [155, 156], and TuRBO [154]. Surrogate-assisted evolutionary algorithms form a parallel literature [16, 17].

Accuracy is not the objective. Several independent results show that the most accurate surrogate is not reliably the one that produces the best decision. Eggensperger et al. found gradient boosting gave the lowest RMSE while random forests gave the highest rank correlation on all nine datasets, and that tree surrogates reproduced optimizer behaviour where GP surrogates sometimes did not [15, 25]. Williams and Cremaschi found MARS and GP best for approximation but random forests, SVR and GP most robust for surrogate-based optimization [28]. Loshchilov et al. use a surrogate that only ranks candidates, because comparison-based evolution strategies need nothing else [247]. Neufang et al. observed optimizer rankings change between test functions and chemical-process cases [29], and Lütjens et al. found the ranking of climate emulators flipped with the noise level of the training data [96]. Theory explains why (Section 9): the regret of the surrogate's optimum is bounded by the surrogate error on the near-optimal sublevel sets, not by its global mean-squared error, and gradient-based optimizers additionally depend on derivative accuracy.

Benchmarks. Test-function libraries—the Surjanovic–Bingham virtual library used here, the 175-function survey of Jamil and Yang, and the COCO/BBOB platform [248, 249]—classify functions by shape, modality, separability and conditioning; exploratory landscape analysis computes such features from samples [250, 251] and supports automated algorithm selection. Benchmarking

practice—performance and data profiles, multiple-problem statistics and their Bayesian alternatives [252–256]—and surrogate-based benchmark suites [257] complete the toolkit.

9 What theory predicts

No free lunch, but structured lunches. Averaged over all functions, no optimizer or learner beats another [258]; every advantage comes from matching an inductive bias to a function class. Information-based complexity makes the price of generality explicit: for functions of smoothness s in d dimensions, the minimax error from n samples scales as $n^{-s/d}$ for approximation and $n^{-s/(2s+d)}$ for noisy regression [259, 260]. Smoothness and dimension, not the method's name, set the achievable rate.

Kernels and GPs. For kernel interpolation the error is controlled by the fill distance h of the design through the power function, at a rate set by the kernel's smoothness, with a trade-off between accuracy and conditioning [43, 44]; GP posterior contraction rates match when the prior's smoothness matches the truth's and degrade otherwise [261]. The fill-distance dependence explains why a spacing statistic of the design predicted accuracy in a large empirical comparison [104], and why space-filling designs are hard to beat for global fit. Expected improvement converges at known rates for functions in the kernel's reproducing-kernel Hilbert space [262], and confidence-bound acquisition has regret bounds [149]—guarantees that hold only when the function is as smooth as the kernel assumes.

Trees and networks. A piecewise-constant partition of n cells approximates a Lipschitz function at rate $n^{-1/d}$, slower than a well-matched kernel on smooth functions but indifferent to discontinuities, plateaus and floors; ensembles are consistent under mild conditions [66]. Networks escape the curse of dimension for functions with suitably bounded spectral norms [263] and approximate smooth functions efficiently with depth [264], but trained networks are biased toward low frequencies [105], and rotation-invariant learners pay linearly in the number of irrelevant inputs [106].

Noise. With observation noise of variance σ^2, any held-out error measured against noisy targets is bounded below by σ^2; differences between models shrink toward zero as σ grows, which is how a model-ranking study can be dominated by its noise level [96]. Clean-truth scoring on synthetic problems removes this floor.

From surrogate error to design quality. If $\hat{x}$ minimises the surrogate s and $x^\star$ minimises f, then $f(\hat{x}) - f(x^\star) \le 2 \sup |f - s|$ over the union of the near-optimal sublevel sets of f and s. Global mean-squared error neither bounds nor is bounded by this quantity; rank fidelity on the sublevel sets governs comparison-based optimizers, and derivative fidelity near $x^\star$ governs gradient-based ones. Multi-fidelity and compression methods enter the same bound through their effect on $\sup |f - s|$ per unit cost.

10 Computational complexity

Asymptotic cost decides which methods are *admissible* at a given sample size, dimension and simulation cost, before accuracy is considered. Table 1 collects fit, storage and prediction costs for the families reviewed above, in the usual algorithm for each; Table 2 does the same for the

representations used to make a learner optimizer-ready. Constants matter at engineering sample sizes, so the study also reports measured times.

Three walls recur. Exact GPs and dense RBF systems need $O(n^3)$ time and $O(n^2)$ memory [30, 43]—about 0.8 GB of kernel matrix at $n = 10^4$—which inducing-point, Nyström and random-feature approximations reduce to $O(nq^2)$ [36, 205, 206], and partitioned or local GPs reduce by working on subsets [37, 39]. Full tensor grids, and hence tensor-product spline upscaling, cost $O(m^d)$: 24 points per axis is 1.4×10^4 nodes in three dimensions, 1.9×10^8 in six and 6×10^{13} in ten; sparse grids [49] and tensor-train formats [201, 202] restore near-linear scaling in d under mixed smoothness or low rank. Delaunay-based interpolants (piecewise-linear, natural neighbour) grow as $O(n^{\lceil d/2 \rceil})$ in the worst case and are impractical beyond about six inputs. By contrast, tree ensembles fit in $O(n \log n)$ per tree and predict in time independent of n [65, 68, 69], and in-context tabular models replace training by $O(n^2)$ attention [81].

Loops multiply these costs. An active-learning loop that refits an exact GP after each of N acquisitions costs $O(N^4)$ unless the Cholesky factor is updated, and joint co-kriging over L fidelities costs $O((\sum_\ell n_\ell)^3)$ while the recursive formulation costs $O(\sum_\ell n_\ell^3)$ [146]. Randomised linear algebra [182, 183] reduces the cost of a decomposition to its intrinsic rank, which is how compressed DMD accelerates its intermediate step. For finite-element simulators costing an hour per run, every surrogate term is negligible up to $n \sim 10^4$; for cheap simulators these terms decide the winner.

Table 1: Asymptotic costs of surrogate learners. n training points, d inputs, P basis terms, s active terms, q inducing points or features, B trees, T boosting rounds, D depth, W weights, E epochs, K ensemble members, HR hyperparameter iterations $\times$ restarts.

Method	Fit time	Memory	Predict (mean / variance)
Polynomial RSM, $P = \binom{d+p}{p}$	$O(nP^2 + P^3)$	$O(nP)$	$O(P)$ / $O(P^2)$
Sparse regression, sparse PCE (OMP)	$O(snP)$	$O(nP)$	$O(s)$
Exact GP / kriging	$O(HR\,n^3)$	$O(n^2)$	$O(nd)$ / $O(n^2)$
Sparse GP, Nyström, random features	$O(nq^2)$	$O(nq)$	$O(q)$ / $O(q^2)$
Treed GP (leaves n_j)	$O(S\sum_j n_j^3)$	$O(\sum_j n_j^2)$	$O(n_j)$
RBF interpolation (dense)	$O(n^3)$	$O(n^2)$	$O(nd)$
IDW / k-NN	$O(dn \log n)$ (tree build)	$O(nd)$	$O(\log n)$–$O(nd)$
Delaunay linear / natural neighbour	$O(n^{\lceil d/2 \rceil})$ worst case	same	$O(\log n)$
MARS, M terms	$O(ndM^2)$	$O(nM)$	$O(M)$
Random forest	$O(Bd'n \log n)$	$O(Bn)$	$O(BD)$
Gradient boosting (histogram)	$O(Tnd)$	$O(nd + T2^D)$	$O(TD)$
BART, m_t trees, S sweeps	$O(Sm_tn)$	$O(m_tn)$	$O(Sm_tD)$
MLP / deep ensemble	$O(KEnW)$	$O(KW)$	$O(KW)$
Deep kernel learning	$O(E(n^3 + nW))$	$O(n^2 + W)$	$O(nd + W)$ / $O(n^2)$
TabPFN (in-context)	—	$O(n^2)$	batched $O(nM)$
Tensor train (cross)	$O(dmr^3)$; $O(dmr^2)$ samples	$O(dmr^2)$	$O(dr^2)$
Compressed sensing (OMP), basis N	$O(snN)$	$O(nN)$	$O(s)$

11 Gap and contributions

The reviewed literature establishes that no surrogate family dominates, that winners are predictable from problem features, that within-family configuration matters, and that approximation accuracy is a weak proxy for decision quality. It does not, to our knowledge, contain a study that jointly varies (i) surface class and landscape features, (ii) dimension and sample budget, (iii) noise level and type, (iv) fidelity structure and cost ratio, (v) the pairing of the sampling scout with the deliverable

Table 2: Costs of optimizer-ready representations; m nodes per axis, spline order k ($k = 4$ cubic), level λ, rank r.

Representation	Build	Storage	Evaluate
Tensor B-spline on dense grid	m^d model calls $+ O(dm^d)$	$O(m^d)$	$O(k^d)$
Multilevel B-spline (MBA)	$O(h(n + m_h^d))$	$O(m_h^d)$	$O(k^d)$
Smolyak sparse grid	$O(2^\lambda \lambda^{d-1})$ model calls	same	$O(2^\lambda \lambda^{d-1})$
Additive P-spline / GAM	$O(ndK)$	$O(dK)$	$O(dk)$
Tensor-train compressed grid	$O(dmr^3)$	$O(dmr^2)$	$O(dr^2)$

model, (vi) the upscaling from learner to optimizer-ready surrogate, including compression, and (vii) the optimizer—and scores complete pipelines by true-function regret per unit simulation cost, while explaining the persistent disagreement between engineering practice (Gaussian processes) and tabular machine learning (tree ensembles). The study part of this paper does this on the Surjanovic–Bingham functions, the multi-fidelity analytical benchmarks, and two real engineering datasets: coupled finite-element sizing of an inflatable ramjet duct, which carries a closed-form low-fidelity model at every point, and historical aircraft performance data.

References

[1] Nestor V. Queipo, Raphael T. Haftka, Wei Shyy, Tushar Goel, Rajkumar Vaidyanathan, and P. Kevin Tucker. Surrogate-based analysis and optimization. *Progress in Aerospace Sciences*, 41(1):1–28, January 2005. ISSN 0376-0421. doi: 10.1016/j.paerosci.2005.02.001. URL http://dx.doi.org/10.1016/j.paerosci.2005.02.001.

[2] Andy Keane, Alexander Forrester, and Andras Sobester. *Engineering Design via Surrogate Modelling: A Practical Guide*. American Institute of Aeronautics and Astronautics, Inc., September 2008. ISBN 9781563479557. doi: 10.2514/4.479557. URL http://dx.doi.org/10.2514/4.479557.

[3] Alexander I.J. Forrester and Andy J. Keane. Recent advances in surrogate-based optimization. *Progress in Aerospace Sciences*, 45(1-3):50–79, January 2009. ISSN 0376-0421. doi: 10.1016/j.paerosci.2008.11.001. URL http://dx.doi.org/10.1016/j.paerosci.2008.11.001.

[4] G. E. P. Box and K. B. Wilson. On the experimental attainment of optimum conditions. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, 13(1):1–38, January 1951. ISSN 1467-9868. doi: 10.1111/j.2517-6161.1951.tb00067.x. URL http://dx.doi.org/10.1111/j.2517-6161.1951.tb00067.x.

[5] Daniel G. Krige. A statistical approach to some basic mine valuation problems on the Witwatersrand. *Journal of the Chemical, Metallurgical and Mining Society of South Africa*, 52 (6):119–139, 1951.

[6] Georges Matheron. Principles of geostatistics. *Economic Geology*, 58(8):1246–1266, December 1963. ISSN 0361-0128. doi: 10.2113/gsecongeo.58.8.1246. URL http://dx.doi.org/10.2113/gsecongeo.58.8.1246.

[7] Noel Cressie. The origins of kriging. *Mathematical Geology*, 22(3):239–252, April 1990. ISSN 1573-8868. doi: 10.1007/bf00889887. URL http://dx.doi.org/10.1007/bf00889887.

[8] Jerome Sacks, William J. Welch, Toby J. Mitchell, and Henry P. Wynn. Design and analysis of computer experiments. *Statistical Science*, 4(4), November 1989. ISSN 0883-4237. doi: 10.1214/ss/1177012413. URL http://dx.doi.org/10.1214/ss/1177012413.

[9] Thomas J. Santner, Brian J. Williams, and William I. Notz. *The Design and Analysis of Computer Experiments*. Springer New York, 2003. ISBN 9781475737998. doi: 10.1007/978-1-4757-3799-8. URL http://dx.doi.org/10.1007/978-1-4757-3799-8.

[10] Peyman Bahrami, Farzan Sahari Moghaddam, and Lesley A. James. A review of proxy modeling highlighting applications for reservoir engineering. *Energies*, 15(14):5247, July 2022. ISSN 1996-1073. doi: 10.3390/en15145247. URL http://dx.doi.org/10.3390/en15145247.

[11] D. Watson-Parris, Y. Rao, D. Olivié, Ø. Seland, P. Nowack, G. Camps-Valls, P. Stier, S. Bouabid, M. Dewey, E. Fons, J. Gonzalez, P. Harder, K. Jeggle, J. Lenhardt, P. Manshausen, M. Novitasari, L. Ricard, and C. Roesch. Climatebench v1.0: A benchmark for data-driven climate projections. *Journal of Advances in Modeling Earth Systems*, 14(10), October 2022. ISSN 1942-2466. doi: 10.1029/2021ms002954. URL http://dx.doi.org/10.1029/2021ms002954.

[12] Katrin Heitmann, David Higdon, Martin White, Salman Habib, Brian J. Williams, Earl Lawrence, and Christian Wagner. The coyote universe. ii. cosmological models and precision emulation of the nonlinear matter power spectrum. *The Astrophysical Journal*, 705(1):156–174, October 2009. ISSN 1538-4357. doi: 10.1088/0004-637x/705/1/156. URL http://dx.doi.org/10.1088/0004-637x/705/1/156.

[13] Ioannis Andrianakis, Ian R. Vernon, Nicky McCreesh, Trevelyan J. McKinley, Jeremy E. Oakley, Rebecca N. Nsubuga, Michael Goldstein, and Richard G. White. Bayesian history matching of complex infectious disease models using emulation: A tutorial and a case study on hiv in uganda. *PLoS Computational Biology*, 11(1):e1003968, January 2015. ISSN 1553-7358. doi: 10.1371/journal.pcbi.1003968. URL http://dx.doi.org/10.1371/journal.pcbi.1003968.

[14] Jack P.C. Kleijnen. *Design and Analysis of Simulation Experiments*. Springer International Publishing, 2015. ISBN 9783319180878. doi: 10.1007/978-3-319-18087-8. URL http://dx.doi.org/10.1007/978-3-319-18087-8.

[15] Katharina Eggensperger, Frank Hutter, Holger Hoos, and Kevin Leyton-Brown. Efficient benchmarking of hyperparameter optimizers via surrogates. *Proceedings of the AAAI Conference on Artificial Intelligence*, 29(1), February 2015. ISSN 2159-5399. doi: 10.1609/aaai.v29i1.9375. URL http://dx.doi.org/10.1609/aaai.v29i1.9375.

[16] Y. Jin. A comprehensive survey of fitness approximation in evolutionary computation. *Soft Computing*, 9(1):3–12, October 2003. ISSN 1433-7479. doi: 10.1007/s00500-003-0328-5. URL http://dx.doi.org/10.1007/s00500-003-0328-5.

[17] Yaochu Jin. Surrogate-assisted evolutionary computation: Recent advances and future challenges. *Swarm and Evolutionary Computation*, 1(2):61–70, June 2011. ISSN 2210-6502. doi: 10.1016/j.swevo.2011.05.001. URL http://dx.doi.org/10.1016/j.swevo.2011.05.001.

[18] Robert B. Gramacy. *Surrogates: Gaussian Process Modeling, Design, and Optimization for the Applied Sciences*. Chapman and Hall/CRC, March 2020. ISBN 9780367815493. doi: 10.1201/9780367815493. URL http://dx.doi.org/10.1201/9780367815493.

[19] M. Kennedy. Predicting the output from a complex computer code when fast approximations are available. *Biometrika*, 87(1):1–13, March 2000. ISSN 1464-3510. doi: 10.1093/biomet/87.1.1. URL http://dx.doi.org/10.1093/biomet/87.1.1.

[20] Alexander I.J Forrester, András Sóbester, and Andy J Keane. Multi-fidelity optimization via surrogate modelling. *Proceedings of the Royal Society A: Mathematical, Physical and Engineering Sciences*, 463(2088):3251–3269, October 2007. ISSN 1471-2946. doi: 10.1098/rspa.2007.1900. URL http://dx.doi.org/10.1098/rspa.2007.1900.

[21] Jacob T. Needels. *Efficient Multidisciplinary Analysis and Optimization of Hypersonic Vehicles Using Multi-Fidelity Surrogate Models*. PhD thesis, Stanford University, 2024. Advisor: J. J. Alonso.

[22] Leo Grinsztajn, Edouard Oyallon, and Gael Varoquaux. Why do tree-based models still outperform deep learning on typical tabular data? In *Advances in Neural Information Processing Systems 35*, NeurIPS 2022, pages 507–520. Neural Information Processing Systems Foundation, Inc. (NeurIPS), 2022. doi: 10.52202/068431-0037. URL http://dx.doi.org/10.52202/068431-0037.

[23] Ravid Shwartz-Ziv and Amitai Armon. Tabular data: Deep learning is not all you need. *Information Fusion*, 81:84–90, May 2022. ISSN 1566-2535. doi: 10.1016/j.inffus.2021.11.011. URL http://dx.doi.org/10.1016/j.inffus.2021.11.011.

[24] Duncan McElfresh, Sujay Khandagale, Jonathan Valverde, Vishak Prasad C, Ganesh Ramakrishnan, Micah Goldblum, and Colin White. When do neural nets outperform boosted trees on tabular data? In *Advances in Neural Information Processing Systems 36*, NeurIPS 2023, pages 76336–76369. Neural Information Processing Systems Foundation, Inc. (NeurIPS), 2023. doi: 10.52202/075280-3337. URL http://dx.doi.org/10.52202/075280-3337.

[25] Katharina Eggensperger, Marius Lindauer, Holger H. Hoos, Frank Hutter, and Kevin Leyton-Brown. Efficient benchmarking of algorithm configurators via model-based surrogates. *Machine Learning*, 107(1):15–41, December 2017. ISSN 1573-0565. doi: 10.1007/s10994-017-5683-z. URL http://dx.doi.org/10.1007/s10994-017-5683-z.

[26] Francesco Lamperti, Andrea Roventini, and Amir Sani. Agent-based model calibration using machine learning surrogates. *Journal of Economic Dynamics and Control*, 90:366–389, May 2018. ISSN 0165-1889. doi: 10.1016/j.jedc.2018.03.011. URL http://dx.doi.org/10.1016/j.jedc.2018.03.011.

[27] Joshua David Poznański. Towards ram air inflated duct eccentric ramjet: A tube launch compatible ramjet. M.s. thesis, University of Kansas, Department of Aerospace Engineering, Lawrence, KS, 2026. Advisor: R. M. Barrett-Gonzalez.

[28] Bianca Williams and Selen Cremaschi. Selection of surrogate modeling techniques for surface approximation and surrogate-based optimization. *Chemical Engineering Research and Design*, 170:76–89, June 2021. ISSN 0263-8762. doi: 10.1016/j.cherd.2021.03.028. URL http://dx.doi.org/10.1016/j.cherd.2021.03.028.

[29] Mathias Neufang, Emma Pajak, Damien van de Berg, Ye Seol Lee, and Ehecatl Antonio del Rio Chanona. Surrogate-based optimization techniques for process systems engineering, 2024.

[30] Carl Edward Rasmussen and Christopher K. I. Williams. *Gaussian Processes for Machine Learning*. The MIT Press, November 2005. ISBN 9780262256834. doi: 10.7551/mitpress/3206.001.0001. URL http://dx.doi.org/10.7551/mitpress/3206.001.0001.

[31] V. Roshan Joseph, Ying Hung, and Agus Sudjianto. Blind kriging: A new method for developing metamodels. *Journal of Mechanical Design*, 130(3), February 2008. ISSN 1528-9001. doi: 10.1115/1.2829873. URL http://dx.doi.org/10.1115/1.2829873.

[32] Max D. Morris, Toby J. Mitchell, and Donald Ylvisaker. Bayesian design and analysis of computer experiments: Use of derivatives in surface prediction. *Technometrics*, 35(3):243–255, August 1993. ISSN 1537-2723. doi: 10.1080/00401706.1993.10485320. URL http://dx.doi.org/10.1080/00401706.1993.10485320.

[33] Mohamed Amine Bouhlel, Nathalie Bartoli, Abdelkader Otsmane, and Joseph Morlier. Improving kriging surrogates of high-dimensional design models by partial least squares dimension reduction. *Structural and Multidisciplinary Optimization*, 53(5):935–952, January 2016. ISSN 1615-1488. doi: 10.1007/s00158-015-1395-9. URL http://dx.doi.org/10.1007/s00158-015-1395-9.

[34] Paul W. Goldberg, Christopher K. I. Williams, and Christopher M. Bishop. Regression with input-dependent noise: A Gaussian process treatment. In *Advances in Neural Information Processing Systems 10*, pages 493–499, 1998.

[35] Bruce Ankenman, Barry L. Nelson, and Jeremy Staum. Stochastic kriging for simulation metamodeling. *Operations Research*, 58(2):371–382, April 2010. ISSN 1526-5463. doi: 10.1287/opre.1090.0754. URL http://dx.doi.org/10.1287/opre.1090.0754.

[36] Michalis Titsias. Variational learning of inducing variables in sparse Gaussian processes. In *Proceedings of the 12th International Conference on Artificial Intelligence and Statistics (AISTATS)*, volume 5 of *PMLR*, pages 567–574, 2009.

[37] Robert B. Gramacy and Daniel W. Apley. Local gaussian process approximation for large computer experiments. *Journal of Computational and Graphical Statistics*, 24(2):561–578, April 2015. ISSN 1537-2715. doi: 10.1080/10618600.2014.914442. URL http://dx.doi.org/10.1080/10618600.2014.914442.

[38] Jasper Snoek, Kevin Swersky, Richard S. Zemel, and Ryan P. Adams. Input warping for bayesian optimization of non-stationary functions, 2014.

[39] Robert B Gramacy and Herbert K. H Lee. Bayesian treed gaussian process models with an application to computer modeling. *Journal of the American Statistical Association*, 103 (483):1119–1130, September 2008. ISSN 1537-274X. doi: 10.1198/016214508000000689. URL http://dx.doi.org/10.1198/016214508000000689.

[40] Andreas C. Damianou and Neil D. Lawrence. Deep gaussian processes, 2012.

[41] Rolland L. Hardy. Multiquadric equations of topography and other irregular surfaces. *Journal of Geophysical Research*, 76(8):1905–1915, March 1971. ISSN 0148-0227. doi: 10.1029/jb076i008p01905. URL http://dx.doi.org/10.1029/jb076i008p01905.

[42] Jean Duchon. *Splines minimizing rotation-invariant semi-norms in Sobolev spaces*, pages 85–100. Springer Berlin Heidelberg, 1977. ISBN 9783540374961. doi: 10.1007/bfb0086566. URL http://dx.doi.org/10.1007/bfb0086566.

[43] Holger Wendland. *Scattered Data Approximation*. Cambridge University Press, December 2004. ISBN 9780511617539. doi: 10.1017/cbo9780511617539. URL http://dx.doi.org/10.1017/cbo9780511617539.

[44] Robert Schaback. Error estimates and condition numbers for radial basis function interpolation. *Advances in Computational Mathematics*, 3(3):251–264, April 1995. ISSN 1572-9044. doi: 10.1007/bf02432002. URL http://dx.doi.org/10.1007/bf02432002.

[45] Gregory E Fasshauer. *Meshfree Approximation Methods with Matlab: (With CD-ROM)*. WORLD SCIENTIFIC, April 2007. ISBN 9789812708632. doi: 10.1142/6437. URL http://dx.doi.org/10.1142/6437.

[46] Holger Wendland. Piecewise polynomial, positive definite and compactly supported radial functions of minimal degree. *Advances in Computational Mathematics*, 4(1):389–396, December 1995. ISSN 1572-9044. doi: 10.1007/bf02123482. URL http://dx.doi.org/10.1007/bf02123482.

[47] Carl de Boor. *A Practical Guide to Splines*. Springer New York, 1978. ISBN 9781461263333. doi: 10.1007/978-1-4612-6333-3. URL http://dx.doi.org/10.1007/978-1-4612-6333-3.

[48] Sergey A. Smolyak. Quadrature and interpolation formulas for tensor products of certain classes of functions. *Soviet Mathematics Doklady*, 4:240–243, 1963.

[49] Hans-Joachim Bungartz and Michael Griebel. Sparse grids. *Acta Numerica*, 13:147–269, May 2004. ISSN 1474-0508. doi: 10.1017/s0962492904000182. URL http://dx.doi.org/10.1017/s0962492904000182.

[50] Lloyd N. Trefethen. *Approximation Theory and Approximation Practice, Extended Edition*. Society for Industrial and Applied Mathematics, January 2019. ISBN 9781611975949. doi: 10.1137/1.9781611975949. URL http://dx.doi.org/10.1137/1.9781611975949.

[51] Yuji Nakatsukasa, Olivier Sète, and Lloyd N. Trefethen. The aaa algorithm for rational approximation. *SIAM Journal on Scientific Computing*, 40(3):A1494–A1522, January 2018. ISSN 1095-7197. doi: 10.1137/16m1106122. URL http://dx.doi.org/10.1137/16m1106122.

[52] Robert Tibshirani. Regression shrinkage and selection via the lasso. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, 58(1):267–288, January 1996. ISSN 1467-9868. doi: 10.1111/j.2517-6161.1996.tb02080.x. URL http://dx.doi.org/10.1111/j.2517-6161.1996.tb02080.x.

[53] Alison Cozad, Nikolaos V. Sahinidis, and David C. Miller. Learning surrogate models for simulation-based optimization. *AIChE Journal*, 60(6):2211–2227, March 2014. ISSN 1547-5905. doi: 10.1002/aic.14418. URL http://dx.doi.org/10.1002/aic.14418.

[54] Steven L. Brunton, Joshua L. Proctor, and J. Nathan Kutz. Discovering governing equations from data by sparse identification of nonlinear dynamical systems. *Proceedings of the National Academy of Sciences*, 113(15):3932–3937, March 2016. ISSN 1091-6490. doi: 10.1073/pnas.1517384113. URL http://dx.doi.org/10.1073/pnas.1517384113.

[55] George S. Kimeldorf and Grace Wahba. A correspondence between bayesian estimation on stochastic processes and smoothing by splines. *The Annals of Mathematical Statistics*, 41

(2):495–502, April 1970. ISSN 0003-4851. doi: 10.1214/aoms/1177697089. URL http://dx.doi.org/10.1214/aoms/1177697089.

[56] Grace Wahba. *Spline Models for Observational Data*. Society for Industrial and Applied Mathematics, January 1990. ISBN 9781611970128. doi: 10.1137/1.9781611970128. URL http://dx.doi.org/10.1137/1.9781611970128.

[57] Paul H. C. Eilers and Brian D. Marx. Flexible smoothing with b-splines and penalties. *Statistical Science*, 11(2), May 1996. ISSN 0883-4237. doi: 10.1214/ss/1038425655. URL http://dx.doi.org/10.1214/ss/1038425655.

[58] Trevor J. Hastie and Robert J. Tibshirani. *Generalized Additive Models*. Chapman and Hall, London, 1990.

[59] Simon N. Wood. Thin plate regression splines. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, 65(1):95–114, January 2003. ISSN 1467-9868. doi: 10.1111/1467-9868.00374. URL http://dx.doi.org/10.1111/1467-9868.00374.

[60] F. N. Fritsch and R. E. Carlson. Monotone piecewise cubic interpolation. *SIAM Journal on Numerical Analysis*, 17(2):238–246, April 1980. ISSN 1095-7170. doi: 10.1137/0717021. URL http://dx.doi.org/10.1137/0717021.

[61] Natalya Pya and Simon N. Wood. Shape constrained additive models. *Statistics and Computing*, 25(3):543–559, February 2014. ISSN 1573-1375. doi: 10.1007/s11222-013-9448-7. URL http://dx.doi.org/10.1007/s11222-013-9448-7.

[62] Jerome H. Friedman. Multivariate adaptive regression splines. *The Annals of Statistics*, 19 (1), March 1991. ISSN 0090-5364. doi: 10.1214/aos/1176347963. URL http://dx.doi.org/10.1214/aos/1176347963.

[63] Jerome H. Friedman and Werner Stuetzle. Projection pursuit regression. *Journal of the American Statistical Association*, 76(376):817–823, December 1981. ISSN 1537-274X. doi: 10.1080/01621459.1981.10477729. URL http://dx.doi.org/10.1080/01621459.1981.10477729.

[64] Leo Breiman, Jerome H. Friedman, Richard A. Olshen, and Charles J. Stone. *Classification and Regression Trees*. Wadsworth, Belmont, CA, 1984.

[65] Leo Breiman. Random forests. *Machine Learning*, 45(1):5–32, October 2001. ISSN 1573-0565. doi: 10.1023/a:1010933404324. URL http://dx.doi.org/10.1023/a:1010933404324.

[66] Gérard Biau and Erwan Scornet. A random forest guided tour. *TEST*, 25(2):197–227, April 2016. ISSN 1863-8260. doi: 10.1007/s11749-016-0481-7. URL http://dx.doi.org/10.1007/s11749-016-0481-7.

[67] Jerome H. Friedman. Greedy function approximation: A gradient boosting machine. *The Annals of Statistics*, 29(5), October 2001. ISSN 0090-5364. doi: 10.1214/aos/1013203451. URL http://dx.doi.org/10.1214/aos/1013203451.

[68] Tianqi Chen and Carlos Guestrin. Xgboost: A scalable tree boosting system. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, KDD '16, pages 785–794. ACM, August 2016. doi: 10.1145/2939672.2939785. URL http://dx.doi.org/10.1145/2939672.2939785.

[69] Guolin Ke, Qi Meng, Thomas Finley, Taifeng Wang, Wei Chen, Weidong Ma, Qiwei Ye, and Tie-Yan Liu. LightGBM: A highly efficient gradient boosting decision tree. In *Advances in Neural Information Processing Systems 30*, pages 3146–3154, 2017.

[70] Liudmila Prokhorenkova, Gleb Gusev, Aleksandr Vorobev, Anna Veronika Dorogush, and Andrey Gulin. Catboost: unbiased boosting with categorical features, 2017.

[71] Hugh A. Chipman, Edward I. George, and Robert E. McCulloch. Bart: Bayesian additive regression trees. *The Annals of Applied Statistics*, 4(1):266–298, March 2010. ISSN 1941-7330. doi: 10.1214/09-aoas285. URL http://dx.doi.org/10.1214/09-aoas285.

[72] Hugh Chipman, Pritam Ranjan, and Weiwei Wang. Sequential design for computer experiments with a flexible bayesian additive model. *Canadian Journal of Statistics*, 40(4):663–678, October 2012. ISSN 1708-945X. doi: 10.1002/cjs.11156. URL http://dx.doi.org/10.1002/cjs.11156.

[73] J. Ross Quinlan. Learning with continuous classes. In *Proceedings of the 5th Australian Joint Conference on Artificial Intelligence*, pages 343–348, 1992.

[74] Adele Cutler. Random forests, September 2014. URL http://dx.doi.org/10.1002/9781118445112.stat06520.

[75] Balaji Lakshminarayanan, Alexander Pritzel, and Charles Blundell. Simple and scalable predictive uncertainty estimation using deep ensembles, 2016.

[76] David J. C. MacKay. A practical bayesian framework for backpropagation networks. *Neural Computation*, 4(3):448–472, May 1992. ISSN 1530-888X. doi: 10.1162/neco.1992.4.3.448. URL http://dx.doi.org/10.1162/neco.1992.4.3.448.

[77] Radford M. Neal. *Bayesian Learning for Neural Networks*. Springer New York, 1996. ISBN 9781461207450. doi: 10.1007/978-1-4612-0745-0. URL http://dx.doi.org/10.1007/978-1-4612-0745-0.

[78] Erik Daxberger, Agustinus Kristiadi, Alexander Immer, Runa Eschenhagen, Matthias Bauer, and Philipp Hennig. Laplace redux – effortless bayesian deep learning, 2021.

[79] Wojciech Marian Czarnecki, Simon Osindero, Max Jaderberg, Grzegorz Świrszcz, and Razvan Pascanu. Sobolev training for neural networks, 2017.

[80] Ziming Liu, Yixuan Wang, Sachin Vaidya, Fabian Ruehle, James Halverson, Marin Soljačić, Thomas Y. Hou, and Max Tegmark. Kan: Kolmogorov-arnold networks, 2024.

[81] Noah Hollmann, Samuel Müller, Lennart Purucker, Arjun Krishnakumar, Max Körfer, Shi Bin Hoo, Robin Tibor Schirrmeister, and Frank Hutter. Accurate predictions on small data with a tabular foundation model. *Nature*, 637(8045):319–326, January 2025. ISSN 1476-4687. doi: 10.1038/s41586-024-08328-6. URL http://dx.doi.org/10.1038/s41586-024-08328-6.

[82] M. Raissi, P. Perdikaris, and G.E. Karniadakis. Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational Physics*, 378:686–707, February 2019. ISSN 0021-9991. doi: 10.1016/j.jcp.2018.10.045. URL http://dx.doi.org/10.1016/j.jcp.2018.10.045.

[83] Lu Lu, Pengzhan Jin, Guofei Pang, Zhongqiang Zhang, and George Em Karniadakis. Learning nonlinear operators via deeponet based on the universal approximation theorem of operators. *Nature Machine Intelligence*, 3(3):218–229, March 2021. ISSN 2522-5839. doi: 10.1038/s42256-021-00302-5. URL http://dx.doi.org/10.1038/s42256-021-00302-5.

[84] Zongyi Li, Nikola Kovachki, Kamyar Azizzadenesheli, Burigede Liu, Kaushik Bhattacharya, Andrew Stuart, and Anima Anandkumar. Fourier neural operator for parametric partial differential equations, 2020.

[85] A. G. Ivakhnenko. Polynomial theory of complex systems. *IEEE Transactions on Systems, Man, and Cybernetics*, SMC-1(4):364–378, October 1971. ISSN 2168-2909. doi: 10.1109/tsmc.1971.4308320. URL http://dx.doi.org/10.1109/tsmc.1971.4308320.

[86] Roger G. Ghanem and Pol D. Spanos. *Stochastic Finite Elements: A Spectral Approach*. Springer New York, 1991. ISBN 9781461230946. doi: 10.1007/978-1-4612-3094-6. URL http://dx.doi.org/10.1007/978-1-4612-3094-6.

[87] Dongbin Xiu and George Em Karniadakis. The Wiener–Askey polynomial chaos for stochastic differential equations. *SIAM Journal on Scientific Computing*, 24(2):619–644, 2002. doi: 10.1137/S1064827501387826.

[88] Nora Lüthen, Stefano Marelli, and Bruno Sudret. Sparse polynomial chaos expansions: Literature survey and benchmark. *SIAM/ASA Journal on Uncertainty Quantification*, 9 (2):593–649, January 2021. ISSN 2166-2525. doi: 10.1137/20m1315774. URL http://dx.doi.org/10.1137/20m1315774.

[89] J.W. Bandler, R.M. Biernacki, Shao Hua Chen, P.A. Grobelny, and R.H. Hemmers. Space mapping technique for electromagnetic optimization. *IEEE Transactions on Microwave Theory and Techniques*, 42(12):2536–2544, 1994. ISSN 0018-9480. doi: 10.1109/22.339794. URL http://dx.doi.org/10.1109/22.339794.

[90] Slawomir Koziel, Qingsha S. Cheng, and John W. Bandler. Space mapping. *IEEE Microwave Magazine*, 9(6):105–122, December 2008. ISSN 1557-9581. doi: 10.1109/mmm.2008.929554. URL http://dx.doi.org/10.1109/mmm.2008.929554.

[91] Donald Shepard. A two-dimensional interpolation function for irregularly-spaced data. In *Proceedings of the 1968 23rd ACM national conference on* -, pages 517–524. ACM Press, 1968. doi: 10.1145/800186.810616. URL http://dx.doi.org/10.1145/800186.810616.

[92] Robin Sibson. A brief description of natural neighbour interpolation. In Vic Barnett, editor, *Interpreting Multivariate Data*, pages 21–36. John Wiley & Sons, 1981.

[93] P. Lancaster and K. Salkauskas. Surfaces generated by moving least squares methods. *Mathematics of Computation*, 37(155):141–158, 1981. ISSN 0025-5718. doi: 10.1090/s0025-5718-1981-0616367-1. URL http://dx.doi.org/10.1090/s0025-5718-1981-0616367-1.

[94] William S. Cleveland. Robust locally weighted regression and smoothing scatterplots. *Journal of the American Statistical Association*, 74(368):829–836, December 1979. ISSN 1537-274X. doi: 10.1080/01621459.1979.10481038. URL http://dx.doi.org/10.1080/01621459.1979.10481038.

[95] Gleb Beliakov. Interpolation of lipschitz functions. *Journal of Computational and Applied Mathematics*, 196(1):20–44, November 2006. ISSN 0377-0427. doi: 10.1016/j.cam.2005.08.011. URL http://dx.doi.org/10.1016/j.cam.2005.08.011.

[96] Björn Lütjens, Raffaele Ferrari, Duncan Watson-Parris, and Noelle E. Selin. The impact of internal variability on benchmarking deep learning climate emulators. *Journal of Advances in Modeling Earth Systems*, 17(8), August 2025. ISSN 1942-2466. doi: 10.1029/2024ms004619. URL http://dx.doi.org/10.1029/2024ms004619.

[97] R. Jin, W. Chen, and T.W. Simpson. Comparative studies of metamodelling techniques under multiple modelling criteria. *Structural and Multidisciplinary Optimization*, 23(1):1–13, December 2001. ISSN 1615-1488. doi: 10.1007/s00158-001-0160-4. URL http://dx.doi.org/10.1007/s00158-001-0160-4.

[98] T.W. Simpson, J.D. Poplinski, P. N. Koch, and J.K. Allen. Metamodels for computer-based engineering design: Survey and recommendations. *Engineering with Computers*, 17(2):129–150, July 2001. ISSN 1435-5663. doi: 10.1007/pl00007198. URL http://dx.doi.org/10.1007/pl00007198.

[99] G. Gary Wang and S. Shan. Review of metamodeling techniques in support of engineering design optimization. *Journal of Mechanical Design*, 129(4):370–380, May 2006. ISSN 1528-9001. doi: 10.1115/1.2429697. URL http://dx.doi.org/10.1115/1.2429697.

[100] Felipe A. C. Viana, Timothy W. Simpson, Vladimir Balabanov, and Vasilli Toropov. Special section on multidisciplinary design optimization: Metamodeling in multidisciplinary design optimization: How far have we really come? *AIAA Journal*, 52(4):670–690, April 2014. ISSN 1533-385X. doi: 10.2514/1.j052375. URL http://dx.doi.org/10.2514/1.j052375.

[101] Reza Alizadeh, Janet K. Allen, and Farrokh Mistree. Managing computational complexity using surrogate models: a critical review. *Research in Engineering Design*, 31(3):275–298, April 2020. ISSN 1435-6066. doi: 10.1007/s00163-020-00336-7. URL http://dx.doi.org/10.1007/s00163-020-00336-7.

[102] Jakub Kudela and Radomil Matousek. Recent advances and applications of surrogate models for finite element method computations: a review. *Soft Computing*, 26(24):13709–13733, July 2022. ISSN 1433-7479. doi: 10.1007/s00500-022-07362-8. URL http://dx.doi.org/10.1007/s00500-022-07362-8.

[103] Raul Yondo, Esther Andrés, and Eusebio Valero. A review on design of experiments and surrogate models in aircraft real-time and many-query aerodynamic analyses. *Progress in Aerospace Sciences*, 96:23–61, January 2018. ISSN 0376-0421. doi: 10.1016/j.paerosci.2017.11.003. URL http://dx.doi.org/10.1016/j.paerosci.2017.11.003.

[104] Bianca Williams and Selen Cremaschi. Novel tool for selecting surrogate modeling techniques for surface approximation. In Metin Türkay and Erdal Aydın, editors, *31st European Symposium on Computer Aided Process Engineering*, volume 50 of *Computer Aided Chemical Engineering*, pages 451–456. Elsevier, 2021. doi: 10.1016/B978-0-323-88506-5.50071-1.

[105] Nasim Rahaman, Aristide Baratin, Devansh Arpit, Felix Draxler, Min Lin, Fred A. Hamprecht, Yoshua Bengio, and Aaron Courville. On the spectral bias of neural networks, 2018.

[106] Andrew Y. Ng. Feature selection, l 1 vs. l 2 regularization, and rotational invariance. In *Twenty-first international conference on Machine learning - ICML '04*, ICML '04, page 78. ACM Press, 2004. doi: 10.1145/1015330.1015435. URL http://dx.doi.org/10.1145/1015330.1015435.

[107] Atharv Bhosekar and Marianthi Ierapetritou. Advances in surrogate based modeling, feasibility analysis, and optimization: A review. *Computers & Chemical Engineering*, 108: 250–267, January 2018. ISSN 0098-1354. doi: 10.1016/j.compchemeng.2017.09.017. URL http://dx.doi.org/10.1016/j.compchemeng.2017.09.017.

[108] Kevin McBride and Kai Sundmacher. Overview of surrogate modeling in chemical process engineering. *Chemie Ingenieur Technik*, 91(3):228–239, January 2019. ISSN 1522-2640. doi: 10.1002/cite.201800091. URL http://dx.doi.org/10.1002/cite.201800091.

[109] Saman Razavi, Bryan A. Tolson, and Donald H. Burn. Review of surrogate modeling in water resources. *Water Resources Research*, 48(7), July 2012. ISSN 1944-7973. doi: 10.1029/2011wr011527. URL http://dx.doi.org/10.1029/2011wr011527.

[110] M. J. Asher, B. F. W. Croke, A. J. Jakeman, and L. J. M. Peeters. A review of surrogate models and their application to groundwater modeling. *Water Resources Research*, 51(8): 5957–5973, August 2015. ISSN 1944-7973. doi: 10.1002/2015wr016967. URL http://dx.doi.org/10.1002/2015wr016967.

[111] Paul Westermann and Ralph Evins. Surrogate modelling for sustainable building design – a review. *Energy and Buildings*, 198:170–186, September 2019. ISSN 0378-7788. doi: 10.1016/j.enbuild.2019.05.057. URL http://dx.doi.org/10.1016/j.enbuild.2019.05.057.

[112] Maliki Moustapha, Stefano Marelli, and Bruno Sudret. Active learning for structural reliability: Survey, general framework and benchmark. *Structural Safety*, 96:102174, May 2022. ISSN 0167-4730. doi: 10.1016/j.strusafe.2021.102174. URL http://dx.doi.org/10.1016/j.strusafe.2021.102174.

[113] M. Ludkovski. Statistical machine learning for quantitative finance. *Annual Review of Statistics and Its Application*, 10(1):271–295, March 2023. ISSN 2326-831X. doi: 10.1146/annurev-statistics-032921-042409. URL http://dx.doi.org/10.1146/annurev-statistics-032921-042409.

[114] Johannes Ruf and Weiguan Wang. Neural networks for option pricing and hedging: A literature review. *SSRN Electronic Journal*, 2019. ISSN 1556-5068. doi: 10.2139/ssrn.3486363. URL http://dx.doi.org/10.2139/ssrn.3486363.

[115] Zhihui Li and Xinqian Zheng. Review of design optimization methods for turbomachinery aerodynamics. *Progress in Aerospace Sciences*, 93:1–23, August 2017. ISSN 0376-0421. doi: 10.1016/j.paerosci.2017.05.003. URL http://dx.doi.org/10.1016/j.paerosci.2017.05.003.

[116] Zhen Hu and Sankaran Mahadevan. Uncertainty quantification and management in additive manufacturing: current status, needs, and opportunities. *The International Journal of Advanced Manufacturing Technology*, 93(5-8):2855–2874, July 2017. ISSN 1433-3015. doi: 10.1007/s00170-017-0703-5. URL http://dx.doi.org/10.1007/s00170-017-0703-5.

[117] Jan N. Fuhg, Amélie Fau, and Udo Nackenhorst. State-of-the-art and comparative review of adaptive sampling methods for kriging. *Archives of Computational Methods in Engineering*, 28(4):2689–2747, August 2020. ISSN 1886-1784. doi: 10.1007/s11831-020-09474-6. URL http://dx.doi.org/10.1007/s11831-020-09474-6.

[118] Tushar Goel, Raphael T. Haftka, Wei Shyy, and Nestor V. Queipo. Ensemble of surrogates. *Structural and Multidisciplinary Optimization*, 33(3):199–216, September 2006. ISSN 1615-1488. doi: 10.1007/s00158-006-0051-9. URL http://dx.doi.org/10.1007/s00158-006-0051-9.

[119] Felipe A. C. Viana, Raphael T. Haftka, and Valder Steffen. Multiple surrogates: how cross-validation errors can help us to obtain the best predictor. *Structural and Multidisciplinary Optimization*, 39(4):439–457, January 2009. ISSN 1615-1488. doi: 10.1007/s00158-008-0338-0. URL http://dx.doi.org/10.1007/s00158-008-0338-0.

[120] David H. Wolpert. Stacked generalization. *Neural Networks*, 5(2):241–259, January 1992. ISSN 0893-6080. doi: 10.1016/s0893-6080(05)80023-1. URL http://dx.doi.org/10.1016/s0893-6080(05)80023-1.

[121] H. S. Seung, M. Opper, and H. Sompolinsky. Query by committee. In *Proceedings of the fifth annual workshop on Computational learning theory*, COLT92, pages 287–294. ACM, July 1992. doi: 10.1145/130385.130417. URL http://dx.doi.org/10.1145/130385.130417.

[122] Stefan Wager, Trevor Hastie, and Bradley Efron. Confidence intervals for random forests: The jackknife and the infinitesimal jackknife. *Journal of Machine Learning Research*, 15: 1625–1651, 2014.

[123] Nicolai Meinshausen. Quantile regression forests. *Journal of Machine Learning Research*, 7: 983–999, 2006.

[124] Tony Duan, Anand Avati, Daisy Yi Ding, Khanh K. Thai, Sanjay Basu, Andrew Y. Ng, and Alejandro Schuler. Ngboost: Natural gradient boosting for probabilistic prediction, 2019.

[125] Andrey Malinin, Liudmila Prokhorenkova, and Aleksei Ustimenko. Uncertainty in gradient boosting via ensembles, 2020.

[126] Yarin Gal and Zoubin Ghahramani. Dropout as a bayesian approximation: Representing model uncertainty in deep learning, 2015.

[127] Alexander Amini, Wilko Schwarting, Ava Soleimany, and Daniela Rus. Deep evidential regression, 2019.

[128] Dennis Ulmer, Christian Hardmeier, and Jes Frellsen. Prior and posterior networks: A survey on evidential deep learning methods for uncertainty estimation, 2021.

[129] Moloud Abdar, Farhad Pourpanah, Sadiq Hussain, Dana Rezazadegan, Li Liu, Mohammad Ghavamzadeh, Paul Fieguth, Xiaochun Cao, Abbas Khosravi, U. Rajendra Acharya, Vladimir Makarenkov, and Saeid Nahavandi. A review of uncertainty quantification in deep learning: Techniques, applications and challenges. *Information Fusion*, 76:243–297, December 2021. ISSN 1566-2535. doi: 10.1016/j.inffus.2021.05.008. URL http://dx.doi.org/10.1016/j.inffus.2021.05.008.

[130] Anastasios N. Angelopoulos and Stephen Bates. *Conformal Prediction: A Gentle Introduction*. Now Publishers, 2023. ISBN 9781638281597. doi: 10.1561/9781638281597. URL http://dx.doi.org/10.1561/9781638281597.

[131] M. D. McKay, R. J. Beckman, and W. J. Conover. A comparison of three methods for selecting values of input variables in the analysis of output from a computer code. *Technometrics*, 21 (2):239, May 1979. ISSN 0040-1706. doi: 10.2307/1268522. URL http://dx.doi.org/10.2307/1268522.

[132] Max D. Morris and Toby J. Mitchell. Exploratory designs for computational experiments. *Journal of Statistical Planning and Inference*, 43(3):381–402, February 1995. ISSN 0378-3758. doi: 10.1016/0378-3758(94)00035-t. URL http://dx.doi.org/10.1016/0378-3758(94)00035-t.

[133] I.M Sobol'. On the distribution of points in a cube and the approximate evaluation of integrals. *USSR Computational Mathematics and Mathematical Physics*, 7(4):86–112, January 1967. ISSN 0041-5553. doi: 10.1016/0041-5553(67)90144-9. URL http://dx.doi.org/10.1016/0041-5553(67)90144-9.

[134] Harald Niederreiter. *Random Number Generation and Quasi-Monte Carlo Methods*. Society for Industrial and Applied Mathematics, January 1992. ISBN 9781611970081. doi: 10.1137/1.9781611970081. URL http://dx.doi.org/10.1137/1.9781611970081.

[135] Qiang Du, Vance Faber, and Max Gunzburger. Centroidal voronoi tessellations: Applications and algorithms. *SIAM Review*, 41(4):637–676, January 1999. ISSN 1095-7200. doi: 10.1137/s0036144599352836. URL http://dx.doi.org/10.1137/s0036144599352836.

[136] Haitao Liu, Yew-Soon Ong, and Jianfei Cai. A survey of adaptive sampling for global metamodeling in support of simulation-based complex engineering design. *Structural and Multidisciplinary Optimization*, 57(1):393–416, June 2017. ISSN 1615-1488. doi: 10.1007/s00158-017-1739-8. URL http://dx.doi.org/10.1007/s00158-017-1739-8.

[137] Sushant S. Garud, Iftekhar A. Karimi, and Markus Kraft. Design of computer experiments: A review. *Computers & Chemical Engineering*, 106:71–95, November 2017. ISSN 0098-1354. doi: 10.1016/j.compchemeng.2017.05.010. URL http://dx.doi.org/10.1016/j.compchemeng.2017.05.010.

[138] Burr Settles. Active learning literature survey. Computer Sciences Technical Report 1648, University of Wisconsin–Madison, 2009.

[139] David J. C. MacKay. Information-based objective functions for active data selection. *Neural Computation*, 4(4):590–604, July 1992. ISSN 1530-888X. doi: 10.1162/neco.1992.4.4.590. URL http://dx.doi.org/10.1162/neco.1992.4.4.590.

[140] David A. Cohn, Zoubin Ghahramani, and Michael I. Jordan. Active learning with statistical models. *Journal of Artificial Intelligence Research*, 4:129–145, 1996. doi: 10.1613/jair.295.

[141] Andreas Krause, Ajit Singh, and Carlos Guestrin. Near-optimal sensor placements in Gaussian processes: Theory, efficient algorithms and empirical studies. *Journal of Machine Learning Research*, 9:235–284, 2008.

[142] Chen Quin Lam. *Sequential Adaptive Designs in Computer Experiments for Response Surface Model Fit*. PhD thesis, The Ohio State University, 2008.

[143] Haitao Liu, Jianfei Cai, and Yew-Soon Ong. An adaptive sampling approach for kriging metamodeling by maximizing expected prediction error. *Computers & Chemical Engineering*, 106:171–182, November 2017. ISSN 0098-1354. doi: 10.1016/j.compchemeng.2017.05.025. URL http://dx.doi.org/10.1016/j.compchemeng.2017.05.025.

[144] Karel Crombecq, Dirk Gorissen, Dirk Deschrijver, and Tom Dhaene. A novel hybrid sequential design strategy for global surrogate modeling of computer experiments. *SIAM Journal on Scientific Computing*, 33(4):1948–1974, January 2011. ISSN 1095-7197. doi: 10.1137/090761811. URL http://dx.doi.org/10.1137/090761811.

[145] Shaoxing Mo, Dan Lu, Xiaoqing Shi, Guannan Zhang, Ming Ye, Jianfeng Wu, and Jichun Wu. A taylor expansion-based adaptive design strategy for global surrogate modeling with applications in groundwater modeling. *Water Resources Research*, 53(12):10802–10823, December 2017. ISSN 1944-7973. doi: 10.1002/2017wr021622. URL http://dx.doi.org/10.1002/2017wr021622.

[146] Loic Le Gratiet and Josselin Garnier. Recursive co-kriging model for design of computer experiments with multiple levels of fidelity. *International Journal for Uncertainty Quantification*, 4 (5):365–386, 2014. ISSN 2152-5080. doi: 10.1615/int.j.uncertaintyquantification.2014006914. URL http://dx.doi.org/10.1615/int.j.uncertaintyquantification.2014006914.

[147] H. J. Kushner. A new method of locating the maximum point of an arbitrary multipeak curve in the presence of noise. *Journal of Basic Engineering*, 86(1):97–106, March 1964. ISSN 0021-9223. doi: 10.1115/1.3653121. URL http://dx.doi.org/10.1115/1.3653121.

[148] Donald R. Jones, Matthias Schonlau, and William J. Welch. Efficient global optimization of expensive black-box functions. *Journal of Global Optimization*, 13(4):455–492, December 1998. ISSN 1573-2916. doi: 10.1023/a:1008306431147. URL http://dx.doi.org/10.1023/a:1008306431147.

[149] Niranjan Srinivas, Andreas Krause, Sham M. Kakade, and Matthias Seeger. Gaussian process optimization in the bandit setting: No regret and experimental design, 2009.

[150] Peter Frazier, Warren Powell, and Savas Dayanik. The knowledge-gradient policy for correlated normal beliefs. *INFORMS Journal on Computing*, 21(4):599–613, November 2009. ISSN 1526-5528. doi: 10.1287/ijoc.1080.0314. URL http://dx.doi.org/10.1287/ijoc.1080.0314.

[151] Philipp Hennig and Christian J. Schuler. Entropy search for information-efficient global optimization, 2011.

[152] Zi Wang and Stefanie Jegelka. Max-value entropy search for efficient Bayesian optimization. In *Proceedings of the 34th International Conference on Machine Learning (ICML)*, volume 70 of *PMLR*, pages 3627–3635, 2017.

[153] Jacob R. Gardner, Matt J. Kusner, Zhixiang Xu, Kilian Q. Weinberger, and John P. Cunningham. Bayesian optimization with inequality constraints. In *Proceedings of the 31st International Conference on Machine Learning (ICML)*, volume 32 of *PMLR*, pages 937–945, 2014.

[154] David Eriksson, Michael Pearce, Jacob R Gardner, Ryan Turner, and Matthias Poloczek. Scalable global optimization via local bayesian optimization, 2019.

[155] Frank Hutter, Holger H. Hoos, and Kevin Leyton-Brown. *Sequential Model-Based Optimization for General Algorithm Configuration*, pages 507–523. Springer Berlin Heidelberg, 2011. ISBN 9783642255663. doi: 10.1007/978-3-642-25566-3_40. URL http://dx.doi.org/10.1007/978-3-642-25566-3_40.

[156] James Bergstra, Rémi Bardenet, Yoshua Bengio, and Balázs Kégl. Algorithms for hyper-parameter optimization. In *Advances in Neural Information Processing Systems 24*, pages 2546–2554, 2011.

[157] Rommel G. Regis and Christine A. Shoemaker. A stochastic radial basis function method for the global optimization of expensive functions. *INFORMS Journal on Computing*, 19 (4):497–509, November 2007. ISSN 1526-5528. doi: 10.1287/ijoc.1060.0182. URL http://dx.doi.org/10.1287/ijoc.1060.0182.

[158] Rommel G. Regis and Christine A. Shoemaker. Combining radial basis function surrogates and dynamic coordinate search in high-dimensional expensive black-box optimization. *Engineering Optimization*, 45(5):529–555, May 2013. ISSN 1029-0273. doi: 10.1080/0305215x.2012.687731. URL http://dx.doi.org/10.1080/0305215x.2012.687731.

[159] B. J. Bichon, M. S. Eldred, L. P. Swiler, S. Mahadevan, and J. M. McFarland. Efficient global reliability analysis for nonlinear implicit performance functions. *AIAA Journal*, 46 (10):2459–2468, October 2008. ISSN 1533-385X. doi: 10.2514/1.34321. URL http://dx.doi.org/10.2514/1.34321.

[160] B. Echard, N. Gayton, and M. Lemaire. Ak-mcs: An active learning reliability method combining kriging and monte carlo simulation. *Structural Safety*, 33(2):145–154, March 2011. ISSN 0167-4730. doi: 10.1016/j.strusafe.2011.01.002. URL http://dx.doi.org/10.1016/j.strusafe.2011.01.002.

[161] Benjamin Peherstorfer, Karen Willcox, and Max Gunzburger. Survey of multifidelity methods in uncertainty propagation, inference, and optimization. *SIAM Review*, 60(3):550–591, January 2018. ISSN 1095-7200. doi: 10.1137/16m1082469. URL http://dx.doi.org/10.1137/16m1082469.

[162] Benjamin Peherstorfer, Karen Willcox, and Max Gunzburger. Optimal model management for multifidelity monte carlo estimation. *SIAM Journal on Scientific Computing*, 38(5):A3163–A3194, January 2016. ISSN 1095-7197. doi: 10.1137/15m1046472. URL http://dx.doi.org/10.1137/15m1046472.

[163] M. Giselle Fernández-Godino. Review of multi-fidelity models, 2016.

[164] Loïc Brevault, Mathieu Balesdent, and Ali Hebbal. Overview of gaussian process based multi-fidelity techniques with variable relationship between fidelities, application to aerospace systems. *Aerospace Science and Technology*, 107:106339, December 2020. ISSN 1270-9638. doi: 10.1016/j.ast.2020.106339. URL http://dx.doi.org/10.1016/j.ast.2020.106339.

[165] Kwan J. Chang, Raphael T. Haftka, Gary L. Giles, and Pi-Jen Kao. Sensitivity-based scaling for approximating structural response. *Journal of Aircraft*, 30(2):283–288, March 1993. ISSN 1533-3868. doi: 10.2514/3.48278. URL http://dx.doi.org/10.2514/3.48278.

[166] Shawn Gano, Brian Sanders, and John Renaud. Variable fidelity optimization using a kriging based scaling function. In *10th AIAA/ISSMO Multidisciplinary Analysis and Optimization*

Conference. American Institute of Aeronautics and Astronautics, August 2004. doi: 10.2514/6.2004-4460. URL http://dx.doi.org/10.2514/6.2004-4460.

[167] N. M. Alexandrov, J. E. Dennis, R. M. Lewis, and V. Torczon. A trust-region framework for managing the use of approximation models in optimization. *Structural Optimization*, 15 (1):16–23, February 1998. ISSN 1615-1488. doi: 10.1007/bf01197433. URL http://dx.doi.org/10.1007/bf01197433.

[168] Natalia M. Alexandrov, Robert Michael Lewis, Clyde R. Gumbert, Lawrence L. Green, and Perry A. Newman. Approximation and model management in aerodynamic optimization with variable-fidelity models. *Journal of Aircraft*, 38(6):1093–1101, November 2001. ISSN 1533-3868. doi: 10.2514/2.2877. URL http://dx.doi.org/10.2514/2.2877.

[169] Zhong-Hua Han and Stefan Görtz. Hierarchical kriging model for variable-fidelity surrogate modeling. *AIAA Journal*, 50(9):1885–1896, September 2012. ISSN 1533-385X. doi: 10.2514/1.j051354. URL http://dx.doi.org/10.2514/1.j051354.

[170] P. Perdikaris, M. Raissi, A. Damianou, N. D. Lawrence, and G. E. Karniadakis. Nonlinear information fusion algorithms for data-efficient multi-fidelity modelling. *Proceedings of the Royal Society A: Mathematical, Physical and Engineering Sciences*, 473(2198): 20160751, February 2017. ISSN 1471-2946. doi: 10.1098/rspa.2016.0751. URL http://dx.doi.org/10.1098/rspa.2016.0751.

[171] Kurt Cutajar, Mark Pullin, Andreas Damianou, Neil Lawrence, and Javier González. Deep gaussian processes for multi-fidelity modeling, 2019.

[172] Xuhui Meng and George Em Karniadakis. A composite neural network that learns from multi-fidelity data: Application to function approximation and inverse pde problems. *Journal of Computational Physics*, 401:109020, January 2020. ISSN 0021-9991. doi: 10.1016/j.jcp.2019.109020. URL http://dx.doi.org/10.1016/j.jcp.2019.109020.

[173] David J. J. Toal. Some considerations regarding the use of multi-fidelity kriging in the construction of surrogate models. *Structural and Multidisciplinary Optimization*, 51(6):1223–1245, December 2014. ISSN 1615-1488. doi: 10.1007/s00158-014-1209-5. URL http://dx.doi.org/10.1007/s00158-014-1209-5.

[174] Sander van Rijn, Sebastian Schmitt, Matthijs van Leeuwen, and Thomas Bäck. Finding efficient trade-offs in multi-fidelity response surface modeling, 2021.

[175] Sander van Rijn and Sebastian Schmitt. Mf2: A collection of multi-fidelity benchmark functions in python. *Journal of Open Source Software*, 5(52):2049, August 2020. ISSN 2475-9066. doi: 10.21105/joss.02049. URL http://dx.doi.org/10.21105/joss.02049.

[176] L. Mainini, A. Serani, M. P. Rumpfkeil, E. Minisci, D. Quagliarella, H. Pehlivan, S. Yildiz, S. Ficini, R. Pellegrini, F. Di Fiore, D. Bryson, M. Nikbay, M. Diez, and P. Beran. Analytical benchmark problems for multifidelity optimization methods, 2022.

[177] Laura Mainini, Andrea Serani, Hayriye Pehlivan-Solak, Francesco Di Fiore, Markus P. Rumpfkeil, Edmondo Minisci, Domenico Quagliarella, Sihmehmet Yildiz, Simone Ficini, Riccardo Pellegrini, Andrew Thelen, Dean Bryson, Melike Nikbay, Matteo Diez, and Philip S. Beran. Analytical benchmark problems and methodological framework for the

assessment and comparison of multifidelity optimization methods. *Archives of Computational Methods in Engineering*, 33(2):2969–3000, November 2025. ISSN 1886-1784. doi: 10.1007/s11831-025-10392-8. URL http://dx.doi.org/10.1007/s11831-025-10392-8.

[178] Kirthevasan Kandasamy, Gautam Dasarathy, Junier B. Oliva, Jeff Schneider, and Barnabás Póczos. Gaussian process bandit optimisation with multi-fidelity evaluations. In *Advances in Neural Information Processing Systems 29*, pages 992–1000, 2016.

[179] Jacob T. Needels and Juan J. Alonso. Trajectory-informed sampling for efficient construction of multi-fidelity surrogate models for hypersonic vehicles. In *AIAA SCITECH 2024 Forum*. American Institute of Aeronautics and Astronautics, January 2024. doi: 10.2514/6.2024-1013. URL http://dx.doi.org/10.2514/6.2024-1013.

[180] Kathleen Champion, Bethany Lusch, J. Nathan Kutz, and Steven L. Brunton. Data-driven discovery of coordinates and governing equations. *Proceedings of the National Academy of Sciences*, 116(45):22445–22451, October 2019. ISSN 1091-6490. doi: 10.1073/pnas.1906995116. URL http://dx.doi.org/10.1073/pnas.1906995116.

[181] N. Benjamin Erichson, Steven L. Brunton, and J. Nathan Kutz. Compressed dynamic mode decomposition for background modeling. *Journal of Real-Time Image Processing*, 16(5): 1479–1492, November 2016. ISSN 1861-8219. doi: 10.1007/s11554-016-0655-2. URL http://dx.doi.org/10.1007/s11554-016-0655-2.

[182] N. Benjamin Erichson, Lionel Mathelin, J. Nathan Kutz, and Steven L. Brunton. Randomized dynamic mode decomposition. *SIAM Journal on Applied Dynamical Systems*, 18(4):1867–1891, January 2019. ISSN 1536-0040. doi: 10.1137/18m1215013. URL http://dx.doi.org/10.1137/18m1215013.

[183] N. Halko, P. G. Martinsson, and J. A. Tropp. Finding structure with randomness: Probabilistic algorithms for constructing approximate matrix decompositions. *SIAM Review*, 53(2):217–288, January 2011. ISSN 1095-7200. doi: 10.1137/090771806. URL http://dx.doi.org/10.1137/090771806.

[184] Paul G. Constantine. *Active Subspaces: Emerging Ideas for Dimension Reduction in Parameter Studies*. Society for Industrial and Applied Mathematics, March 2015. ISBN 9781611973860. doi: 10.1137/1.9781611973860. URL http://dx.doi.org/10.1137/1.9781611973860.

[185] Jeffrey M. Hokanson and Paul G. Constantine. Data-driven polynomial ridge approximation using variable projection. *SIAM Journal on Scientific Computing*, 40(3):A1566–A1589, January 2018. ISSN 1095-7197. doi: 10.1137/17m1117690. URL http://dx.doi.org/10.1137/17m1117690.

[186] Rohit Tripathy, Ilias Bilionis, and Marcial Gonzalez. Gaussian processes with built-in dimensionality reduction: Applications to high-dimensional uncertainty propagation. *Journal of Computational Physics*, 321:191–223, September 2016. ISSN 0021-9991. doi: 10.1016/j.jcp.2016.05.039. URL http://dx.doi.org/10.1016/j.jcp.2016.05.039.

[187] Ziyu Wang, Frank Hutter, Masrour Zoghi, David Matheson, and Nando De Feitas. Bayesian optimization in a billion dimensions via random embeddings. *Journal of Artificial Intelligence Research*, 55:361–387, February 2016. ISSN 1076-9757. doi: 10.1613/jair.4806. URL http://dx.doi.org/10.1613/jair.4806.

[188] Andrew Gordon Wilson, Zhiting Hu, Ruslan Salakhutdinov, and Eric P. Xing. Deep kernel learning, 2015.

[189] Sebastian W. Ober, Carl E. Rasmussen, and Mark van der Wilk. The promises and pitfalls of deep kernel learning, 2021.

[190] Houman Owhadi and Gene Ryan Yoo. Kernel flows: From learning kernels from data into the abyss. *Journal of Computational Physics*, 389:22–47, July 2019. ISSN 0021-9991. doi: 10.1016/j.jcp.2019.03.040. URL http://dx.doi.org/10.1016/j.jcp.2019.03.040.

[191] Matthew Tancik, Pratul P. Srinivasan, Ben Mildenhall, Sara Fridovich-Keil, Nithin Raghavan, Utkarsh Singhal, Ravi Ramamoorthi, Jonathan T. Barron, and Ren Ng. Fourier features let networks learn high frequency functions in low dimensional domains, 2020.

[192] Andrew Gordon Wilson and Ryan Prescott Adams. Gaussian process kernels for pattern discovery and extrapolation, 2013.

[193] Rafael Gómez-Bombarelli, Jennifer N. Wei, David Duvenaud, José Miguel Hernández-Lobato, Benjamín Sánchez-Lengeling, Dennis Sheberla, Jorge Aguilera-Iparraguirre, Timothy D. Hirzel, Ryan P. Adams, and Alán Aspuru-Guzik. Automatic chemical design using a data-driven continuous representation of molecules. *ACS Central Science*, 4(2): 268–276, January 2018. ISSN 2374-7951. doi: 10.1021/acscentsci.7b00572. URL http://dx.doi.org/10.1021/acscentsci.7b00572.

[194] E.J. Candes, J. Romberg, and T. Tao. Robust uncertainty principles: exact signal reconstruction from highly incomplete frequency information. *IEEE Transactions on Information Theory*, 52(2):489–509, February 2006. ISSN 0018-9448. doi: 10.1109/tit.2005.862083. URL http://dx.doi.org/10.1109/tit.2005.862083.

[195] D.L. Donoho. Compressed sensing. *IEEE Transactions on Information Theory*, 52(4):1289–1306, April 2006. ISSN 0018-9448. doi: 10.1109/tit.2006.871582. URL http://dx.doi.org/10.1109/tit.2006.871582.

[196] Holger Rauhut and Rachel Ward. Sparse legendre expansions via $\ell 1$-minimization. *Journal of Approximation Theory*, 164(5):517–533, May 2012. ISSN 0021-9045. doi: 10.1016/j.jat.2012.01.008. URL http://dx.doi.org/10.1016/j.jat.2012.01.008.

[197] Alireza Doostan and Houman Owhadi. A non-adapted sparse approximation of pdes with stochastic inputs. *Journal of Computational Physics*, 230(8):3015–3034, April 2011. ISSN 0021-9991. doi: 10.1016/j.jcp.2011.01.002. URL http://dx.doi.org/10.1016/j.jcp.2011.01.002.

[198] Jerrad Hampton and Alireza Doostan. Compressive sampling of polynomial chaos expansions: Convergence analysis and sampling strategies. *Journal of Computational Physics*, 280:363–386, January 2015. ISSN 0021-9991. doi: 10.1016/j.jcp.2014.09.019. URL http://dx.doi.org/10.1016/j.jcp.2014.09.019.

[199] David L Donoho and Iain M Johnstone. Ideal spatial adaptation by wavelet shrinkage. *Biometrika*, 81(3):425–455, September 1994. ISSN 1464-3510. doi: 10.1093/biomet/81.3.425. URL http://dx.doi.org/10.1093/biomet/81.3.425.

[200] Alex Townsend and Lloyd N. Trefethen. An extension of chebfun to two dimensions. *SIAM Journal on Scientific Computing*, 35(6):C495–C518, January 2013. ISSN 1095-7197. doi: 10.1137/130908002. URL http://dx.doi.org/10.1137/130908002.

[201] I. V. Oseledets. Tensor-train decomposition. *SIAM Journal on Scientific Computing*, 33(5): 2295–2317, January 2011. ISSN 1095-7197. doi: 10.1137/090752286. URL http://dx.doi.org/10.1137/090752286.

[202] Ivan Oseledets and Eugene Tyrtyshnikov. Tt-cross approximation for multidimensional arrays. *Linear Algebra and its Applications*, 432(1):70–88, January 2010. ISSN 0024-3795. doi: 10.1016/j.laa.2009.07.024. URL http://dx.doi.org/10.1016/j.laa.2009.07.024.

[203] Alex Gorodetsky, Sertac Karaman, and Youssef Marzouk. A continuous analogue of the tensor-train decomposition. *Computer Methods in Applied Mechanics and Engineering*, 347: 59–84, April 2019. ISSN 0045-7825. doi: 10.1016/j.cma.2018.12.015. URL http://dx.doi.org/10.1016/j.cma.2018.12.015.

[204] Rachel Minster, Arvind K. Saibaba, and Misha E. Kilmer. Randomized algorithms for low-rank tensor decompositions in the tucker format. *SIAM Journal on Mathematics of Data Science*, 2(1):189–215, January 2020. ISSN 2577-0187. doi: 10.1137/19m1261043. URL http://dx.doi.org/10.1137/19m1261043.

[205] Christopher K. I. Williams and Matthias Seeger. Using the Nyström method to speed up kernel machines. In *Advances in Neural Information Processing Systems 13*, pages 682–688, 2001.

[206] Ali Rahimi and Benjamin Recht. Random features for large-scale kernel machines. In *Advances in Neural Information Processing Systems 20*, pages 1177–1184, 2007.

[207] David P. Woodruff. Sketching as a tool for numerical linear algebra. *Foundations and Trends® in Theoretical Computer Science*, 10(1-2):1–157, October 2014. ISSN 1551-3068. doi: 10.1561/0400000060. URL http://dx.doi.org/10.1561/0400000060.

[208] Chris Thornton, Frank Hutter, Holger H. Hoos, and Kevin Leyton-Brown. Auto-weka: combined selection and hyperparameter optimization of classification algorithms. In *Proceedings of the 19th ACM SIGKDD international conference on Knowledge discovery and data mining*, KDD' 13, pages 847–855. ACM, August 2013. doi: 10.1145/2487575.2487629. URL http://dx.doi.org/10.1145/2487575.2487629.

[209] Matthias Feurer, Aaron Klein, Katharina Eggensperger, Jost Springenberg, Manuel Blum, and Frank Hutter. Efficient and robust automated machine learning. In *Advances in Neural Information Processing Systems 28*, pages 2962–2970, 2015.

[210] Andreas Wächter and Lorenz T. Biegler. On the implementation of an interior-point filter line-search algorithm for large-scale nonlinear programming. *Mathematical Programming*, 106(1):25–57, April 2005. ISSN 1436-4646. doi: 10.1007/s10107-004-0559-y. URL http://dx.doi.org/10.1007/s10107-004-0559-y.

[211] S. Lee, G. Wolberg, and S.Y. Shin. Scattered data interpolation with multilevel b-splines. *IEEE Transactions on Visualization and Computer Graphics*, 3(3):228–244, 1997. ISSN 1077-2626. doi: 10.1109/2945.620490. URL http://dx.doi.org/10.1109/2945.620490.

[212] David R. Forsey and Richard H. Bartels. Hierarchical b-spline refinement. *ACM SIGGRAPH Computer Graphics*, 22(4):205–212, June 1988. ISSN 0097-8930. doi: 10.1145/378456.378512. URL http://dx.doi.org/10.1145/378456.378512.

[213] Carlotta Giannelli, Bert Jüttler, and Hendrik Speleers. Thb-splines: The truncated basis for hierarchical splines. *Computer Aided Geometric Design*, 29(7):485–498, October 2012. ISSN 0167-8396. doi: 10.1016/j.cagd.2012.03.025. URL http://dx.doi.org/10.1016/j.cagd.2012.03.025.

[214] Erdem Acar and Masoud Rais-Rohani. Ensemble of metamodels with optimized weight factors. *Structural and Multidisciplinary Optimization*, 37(3):279–294, 2009. doi: 10.1007/s00158-008-0230-y.

[215] Hao Chen, Zhilang Zhang, Weikun Li, Qimeng Liu, Kai Sun, Dixia Fan, and Weicheng Cui. Ensemble of surrogates in black-box-type engineering optimization: Recent advances and applications. *Expert Systems with Applications*, 248:123427, 2024.

[216] Leo Breiman. Bagging predictors. *Machine Learning*, 24(2):123–140, 1996. doi: 10.1007/BF00058655.

[217] Peter Bühlmann and Bin Yu. Analyzing bagging. *The Annals of Statistics*, 30(4):927–961, 2002.

[218] Ian Osband, John Aslanides, and Albin Cassirer. Randomized prior functions for deep reinforcement learning. In *Advances in Neural Information Processing Systems 31 (NeurIPS)*, 2018.

[219] Felipe A. C. Viana, Raphael T. Haftka, and Layne T. Watson. Efficient global optimization algorithm assisted by multiple surrogate techniques. *Journal of Global Optimization*, 56(2):669–689, 2013. doi: 10.1007/s10898-012-9892-5.

[220] Jack P. C. Kleijnen, Wim van Beers, and Inneke van Nieuwenhuyse. Expected improvement in efficient global optimization through bootstrapped kriging. *Journal of Global Optimization*, 54(1):59–73, 2012. doi: 10.1007/s10898-011-9741-y.

[221] Qin Lu, Konstantinos D. Polyzos, Bingcong Li, and Georgios B. Giannakis. Surrogate modeling for Bayesian optimization beyond a single Gaussian process, 2022.

[222] Zi Wang, Clement Gehring, Pushmeet Kohli, and Stefanie Jegelka. Batched large-scale Bayesian optimization in high-dimensional spaces. In *Proceedings of the 21st International Conference on Artificial Intelligence and Statistics (AISTATS)*, volume 84 of *PMLR*, 2018.

[223] Jost Tobias Springenberg, Aaron Klein, Stefan Falkner, and Frank Hutter. Bayesian optimization with robust Bayesian neural networks. In *Advances in Neural Information Processing Systems 29 (NIPS)*, 2016.

[224] Yucen Lily Li, Tim G. J. Rudner, and Andrew Gordon Wilson. A study of Bayesian neural network surrogates for Bayesian optimization. In *International Conference on Learning Representations (ICLR)*, 2024.

[225] Dudy Lim, Yaochu Jin, Yew-Soon Ong, and Bernhard Sendhoff. Generalizing surrogate-assisted evolutionary computation. *IEEE Transactions on Evolutionary Computation*, 14(3):329–355, 2010. doi: 10.1109/TEVC.2009.2027359.

[226] Maxim Raginsky, Alexander Rakhlin, and Matus Telgarsky. Non-convex learning via stochastic gradient Langevin dynamics: a nonasymptotic analysis. In *Proceedings of the 2017 Conference on Learning Theory (COLT)*, volume 65 of *PMLR*, pages 1674–1703, 2017.

[227] Yurii Nesterov and Vladimir Spokoiny. Random gradient-free minimization of convex functions. *Foundations of Computational Mathematics*, 17(2):527–566, 2017. doi: 10.1007/s10208-015-9296-2.

[228] John C. Duchi, Peter L. Bartlett, and Martin J. Wainwright. Randomized smoothing for stochastic optimization. *SIAM Journal on Optimization*, 22(2):674–701, 2012. doi: 10.1137/110831659.

[229] René Pinnau, Claudia Totzeck, Oliver Tse, and Stephan Martin. A consensus-based model for global optimization and its mean-field limit. *Mathematical Models and Methods in Applied Sciences*, 27(1):183–204, 2017. doi: 10.1142/S0218202517400061.

[230] Massimo Fornasier, Timo Klock, and Konstantin Riedl. Consensus-based optimization methods converge globally. *SIAM Journal on Optimization*, 2024.

[231] Marco A. Iglesias, Kody J. H. Law, and Andrew M. Stuart. Ensemble Kalman methods for inverse problems. *Inverse Problems*, 29(4):045001, 2013. doi: 10.1088/0266-5611/29/4/045001.

[232] Ivo Babuška and Jens M. Melenk. The partition of unity method. *International Journal for Numerical Methods in Engineering*, 40(4):727–758, 1997.

[233] Waltraud Huyer and Arnold Neumaier. Snobfit – stable noisy optimization by branch and fit. *ACM Transactions on Mathematical Software*, 35(2):1–25, July 2008. ISSN 1557-7295. doi: 10.1145/1377612.1377613. URL http://dx.doi.org/10.1145/1377612.1377613.

[234] Velibor V. Mišić. Optimization of tree ensembles. *Operations Research*, 68(5):1605–1624, 2020. doi: 10.1287/opre.2019.1928.

[235] Alexander Thebelt, Jan Kronqvist, Miten Mistry, Robert M. Lee, Nathan Sudermann-Merx, and Ruth Misener. ENTMOOT: A framework for optimization over ensemble tree models. *Computers & Chemical Engineering*, 151:107343, 2021.

[236] Alexander Thebelt, Calvin Tsay, Robert M. Lee, Nathan Sudermann-Merx, David Walz, Behrang Shafei, and Ruth Misener. Tree ensemble kernels for Bayesian optimization with known constraints over mixed-feature spaces. In *Advances in Neural Information Processing Systems 35 (NeurIPS)*, 2022.

[237] Jorge Nocedal and Stephen J. Wright. *Numerical Optimization*. Springer New York, 2006. ISBN 9780387400655. doi: 10.1007/978-0-387-40065-5. URL http://dx.doi.org/10.1007/978-0-387-40065-5.

[238] Dieter Kraft. A software package for sequential quadratic programming. Technical Report DFVLR-FB 88-28, DFVLR, Oberpfaffenhofen, Germany, 1988.

[239] J. A. Nelder and R. Mead. A simplex method for function minimization. *The Computer Journal*, 7(4):308–313, January 1965. ISSN 1460-2067. doi: 10.1093/comjnl/7.4.308. URL http://dx.doi.org/10.1093/comjnl/7.4.308.

[240] M. J. D. Powell. *A Direct Search Optimization Method That Models the Objective and Constraint Functions by Linear Interpolation*, pages 51–67. Springer Netherlands, 1994. ISBN 9789401583305. doi: 10.1007/978-94-015-8330-5_4. URL http://dx.doi.org/10.1007/978-94-015-8330-5_4.

[241] Charles Audet and J. E. Dennis. Mesh adaptive direct search algorithms for constrained optimization. *SIAM Journal on Optimization*, 17(1):188–217, January 2006. ISSN 1095-7189. doi: 10.1137/040603371. URL http://dx.doi.org/10.1137/040603371.

[242] D. R. Jones, C. D. Perttunen, and B. E. Stuckman. Lipschitzian optimization without the lipschitz constant. *Journal of Optimization Theory and Applications*, 79(1):157–181, October 1993. ISSN 1573-2878. doi: 10.1007/bf00941892. URL http://dx.doi.org/10.1007/bf00941892.

[243] Rainer Storn and Kenneth Price. Differential evolution – a simple and efficient heuristic for global optimization over continuous spaces. *Journal of Global Optimization*, 11(4):341–359, December 1997. ISSN 1573-2916. doi: 10.1023/a:1008202821328. URL http://dx.doi.org/10.1023/a:1008202821328.

[244] Nikolaus Hansen and Andreas Ostermeier. Completely derandomized self-adaptation in evolution strategies. *Evolutionary Computation*, 9(2):159–195, June 2001. ISSN 1530-9304. doi: 10.1162/106365601750190398. URL http://dx.doi.org/10.1162/106365601750190398.

[245] James Kennedy and Russell Eberhart. Particle swarm optimization. In *Proceedings of ICNN'95 – International Conference on Neural Networks*, volume 4, pages 1942–1948, 1995. doi: 10.1109/ICNN.1995.488968.

[246] David J. Wales and Jonathan P. K. Doye. Global optimization by basin-hopping and the lowest energy structures of lennard-jones clusters containing up to 110 atoms. *The Journal of Physical Chemistry A*, 101(28):5111–5116, July 1997. ISSN 1520-5215. doi: 10.1021/jp970984n. URL http://dx.doi.org/10.1021/jp970984n.

[247] Ilya Loshchilov, Marc Schoenauer, and Michele Sebag. Self-adaptive surrogate-assisted covariance matrix adaptation evolution strategy. In *Proceedings of the 14th annual conference on Genetic and evolutionary computation*, GECCO '12, pages 321–328. ACM, July 2012. doi: 10.1145/2330163.2330210. URL http://dx.doi.org/10.1145/2330163.2330210.

[248] Momin Jamil and Xin She Yang. A literature survey of benchmark functions for global optimisation problems. *International Journal of Mathematical Modelling and Numerical Optimisation*, 4(2):150, 2013. ISSN 2040-3615. doi: 10.1504/ijmmno.2013.055204. URL http://dx.doi.org/10.1504/ijmmno.2013.055204.

[249] Nikolaus Hansen, Anne Auger, Raymond Ros, Olaf Mersmann, Tea Tušar, and Dimo Brockhoff. Coco: a platform for comparing continuous optimizers in a black-box setting. *Optimization Methods and Software*, 36(1):114–144, August 2020. ISSN 1029-4937. doi: 10.1080/10556788.2020.1808977. URL http://dx.doi.org/10.1080/10556788.2020.1808977.

[250] Olaf Mersmann, Bernd Bischl, Heike Trautmann, Mike Preuss, Claus Weihs, and Günter Rudolph. Exploratory landscape analysis. In *Proceedings of the 13th annual conference on Genetic and evolutionary computation*, GECCO '11, pages 829–836. ACM, July 2011. doi: 10.1145/2001576.2001690. URL http://dx.doi.org/10.1145/2001576.2001690.

[251] Pascal Kerschke and Heike Trautmann. *Comprehensive Feature-Based Landscape Analysis of Continuous and Constrained Optimization Problems Using the R-Package Flacco*, pages 93–123. Springer International Publishing, 2019. ISBN 9783030251475. doi: 10.1007/978-3-030-25147-5_7. URL http://dx.doi.org/10.1007/978-3-030-25147-5_7.

[252] Elizabeth D. Dolan and Jorge J. Moré. Benchmarking optimization software with performance profiles. *Mathematical Programming*, 91(2):201–213, January 2002. ISSN 1436-4646. doi: 10.1007/s101070100263. URL http://dx.doi.org/10.1007/s101070100263.

[253] Jorge J. Moré and Stefan M. Wild. Benchmarking derivative-free optimization algorithms. *SIAM Journal on Optimization*, 20(1):172–191, January 2009. ISSN 1095-7189. doi: 10.1137/080724083. URL http://dx.doi.org/10.1137/080724083.

[254] Janez Demšar. Statistical comparisons of classifiers over multiple data sets. *Journal of Machine Learning Research*, 7:1–30, 2006.

[255] Alessio Benavoli, Giorgio Corani, Janez Demsar, and Marco Zaffalon. Time for a change: a tutorial for comparing multiple classifiers through bayesian analysis, 2016.

[256] Thomas Bartz-Beielstein, Carola Doerr, Daan van den Berg, Jakob Bossek, Sowmya Chandrasekaran, Tome Eftimov, Andreas Fischbach, Pascal Kerschke, William La Cava, Manuel Lopez-Ibanez, Katherine M. Malan, Jason H. Moore, Boris Naujoks, Patryk Orzechowski, Vanessa Volz, Markus Wagner, and Thomas Weise. Benchmarking in optimization: Best practice and open issues, 2020.

[257] Florian Pfisterer, Lennart Schneider, Julia Moosbauer, Martin Binder, and Bernd Bischl. Yahpo gym – an efficient multi-objective multi-fidelity benchmark for hyperparameter optimization, 2021.

[258] D.H. Wolpert and W.G. Macready. No free lunch theorems for optimization. *IEEE Transactions on Evolutionary Computation*, 1(1):67–82, April 1997. ISSN 1089-778X. doi: 10.1109/4235.585893. URL http://dx.doi.org/10.1109/4235.585893.

[259] H. Woźniakowski. *Tractability of Multivariate Problems*, pages 236–276. Cambridge University Press, July 2009. ISBN 9781139107068. doi: 10.1017/cbo9781139107068.009. URL http://dx.doi.org/10.1017/cbo9781139107068.009.

[260] Charles J. Stone. Optimal global rates of convergence for nonparametric regression. *The Annals of Statistics*, 10(4), December 1982. ISSN 0090-5364. doi: 10.1214/aos/1176345969. URL http://dx.doi.org/10.1214/aos/1176345969.

[261] Aad van der Vaart and Harry van Zanten. Information rates of nonparametric Gaussian process methods. *Journal of Machine Learning Research*, 12:2095–2119, 2011.

[262] Adam D. Bull. Convergence rates of efficient global optimization algorithms, 2011.

[263] A.R. Barron. Universal approximation bounds for superpositions of a sigmoidal function. *IEEE Transactions on Information Theory*, 39(3):930–945, May 1993. ISSN 1557-9654. doi: 10.1109/18.256500. URL http://dx.doi.org/10.1109/18.256500.

[264] Dmitry Yarotsky. Error bounds for approximations with deep relu networks. *Neural Networks*, 94:103–114, October 2017. ISSN 0893-6080. doi: 10.1016/j.neunet.2017.07.002. URL http://dx.doi.org/10.1016/j.neunet.2017.07.002.